

Database Compression

Kelly Jiang, Jax Luo

Sqlite

json

↓ spark

SQLite

↓ vertically partition

SQLite

↓

Columnar
Compression

SQLite

↓

Dictionary encoding

↓

Bitpacking

↓

zopf1:

Dictionary Encoding

values

"abc"

"def"

"abc"

"abc"

"efg"



keys

0

1

0

0

2

Dictionary

keys values

0

"abc"

1

"def"

2


"efg"

Bitpacking

Compute Max & Min to figure out bits needed

Truncate the bits representation for each value
in that column

Concatenate bit strings together



0110

0111

1000

→ Output 0110 0111 1000

Discard

String Compression

Zopfli: Made by Google

Supports Deflate, gzip, or zlib encoding

3-8% more compression than zlib max,

~80x slower