

# 01: Introduction

Andrew Crotty // CS497 // Fall 2023

# Today's Agenda

---

Course Logistics

History of Database Systems

Modern OLAP

# Course Objectives

---

Comprehensive overview of the internals of modern OLAP DBMSs.

Students will learn to:

- Read + evaluate systems research papers
- Identify system tradeoffs and justify design decisions
- Craft + deliver presentations to convey research ideas
- Plan + execute a final project that answers an interesting systems research question

This is **not** a course about classical DBMSs.

# Background

---

This course is designed for students interested in systems research (grads + advanced undergrads).

I assume you have already taken an intro DB systems course (e.g., CS339) or equivalent.

**Things that we will not cover:**

Intro DBMS concepts like SQL, Serializability Theory, Relational Algebra, Basic Algorithms + Data Structures, etc.

# Topics

---

Data Storage

Compression

Vectorized Execution

Query Compilation

Parallel Join Algorithms

Index Structures

# Course Logistics

---

## Course Policies + Schedule:

→ Refer to [course web page](#)

## Academic Integrity:

→ Refer to [Northwestern policy page](#)

→ If you're not sure, ask me.

→ Seriously, don't plagiarize or you **will** get wrecked.

# Grading Rubric

---

**Compression Project – 40%**

**Encyclopedia Article – 10%**

**Final Project – 50%**

# Compression Project

---

Groups of 1-2 students will be given 3 real-world datasets to compress as much as possible.

Deliverables include:

- **Project presentation** (5-10 minutes, modeled after a short conference talk) – 5%
- **Written report** (at least 4 pages excluding references, modeled after a workshop paper) – 10%
- **Programming component** (code, benchmarks, demo, etc.) – 25%



# Encyclopedia Article

---

The [Database of Databases](#) is an online encyclopedia of DBMSs maintained by the CMU DB Group.

Groups of 1-2 students will write an article.

- Must provide citations and attributions
- Avoid unscientific (i.e., marketing) language

You may **not** copy text / images directly from papers or other sources.

# Final Project

---

Groups of 1-2 will complete a final project on an approved topic related to the course content.

Deliverables include:

- **Project proposal** (5-10 minutes, modeled after a short conference talk) – 5%
- **Project presentation** (20-30 minutes, modeled after a full conference talk) – 10%
- **Written report** (at least 6 pages excluding references, modeled after a conference paper) – 10%
- **Programming component** (code, benchmarks, demo, etc.) – 25%



# HISTORY OF DATABASE SYSTEMS



WHAT GOES AROUND COMES AROUND  
*READINGS IN DB SYSTEMS, 4TH EDITION, 2006.*



WHAT GOES AROUND COMES AROUND... AND AROUND  
*UNDER SUBMISSION 2023*

# History Repeats Itself

---

Old issues are still relevant today. Many of the ideas in today's systems are not new.

Every decade, someone invents a SQL replacement and some combination of the following happens.

- It fails.
- The project slowly adds SQL features.
- The SQL standard absorbs the best parts.

# 1960s – IDS

---

## Integrated Data Store

Developed internally at GE in early 1960s.

GE sold their computing division to Honeywell in 1969.

One of the first DBMSs.



**Honeywell**

# 1960s – CODASYL

---

COBOL people got together and proposed a standard for how programs will access a database. Lead by [Charles Bachman](#).

- Network data model
- Tuple-at-a-time queries

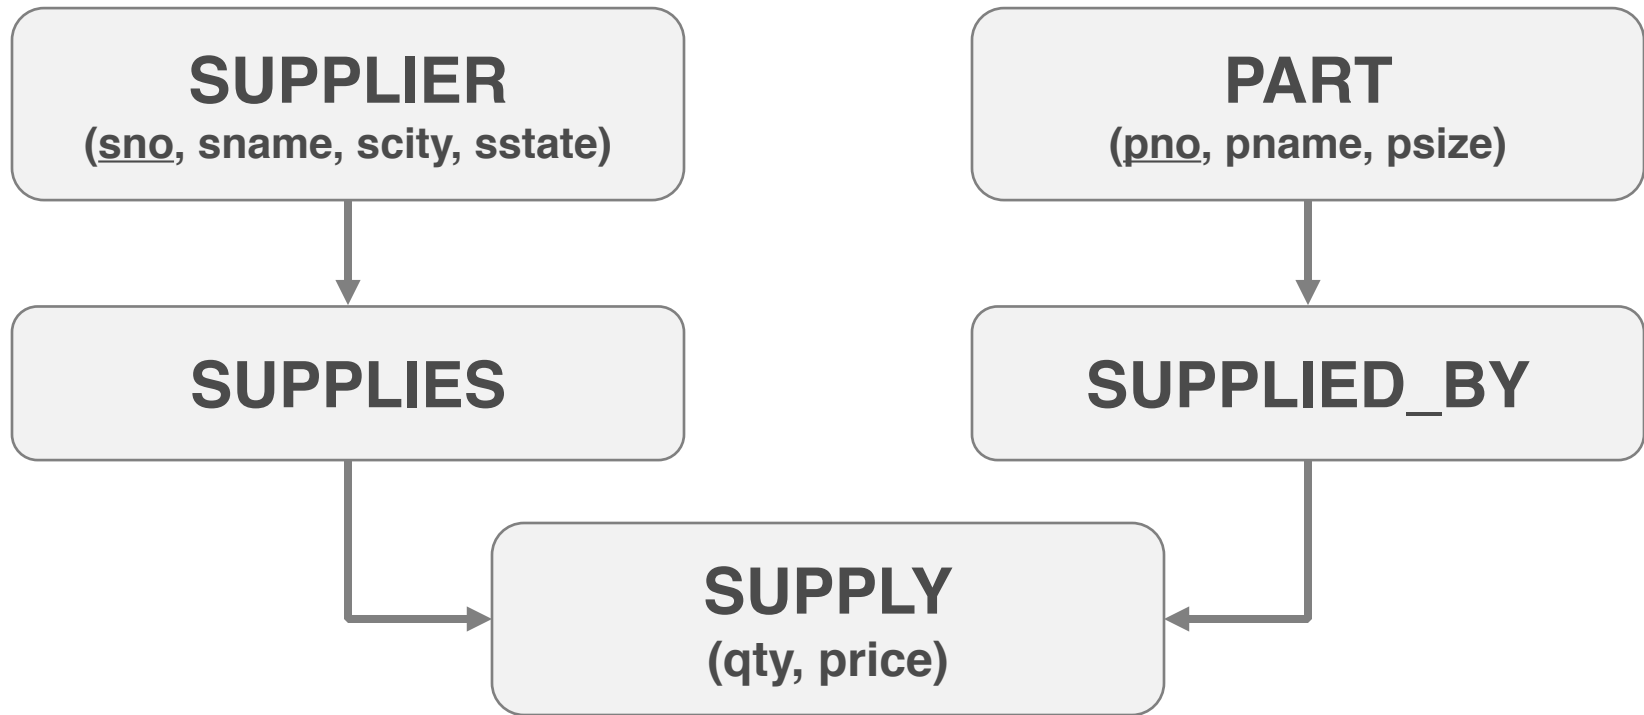


Bachman

Bachman also worked at [Culliane Database Systems](#) in the 1970s to help build **IDMS**.

# Network Data Model

## *Schema*



# Network Data Model

## *Instance*

### SUPPLIER

sno	sname	scity	sstate
1001	Dirty Rick	New York	NY
1002	Squirrels	Boston	MA

### PART

pno	pname	psize
999	Batteries	Large

### SUPPLY

qty	price
10	\$100
14	\$99



# Network Data Model

## *Instance*

### SUPPLIER

sno	sname	scity	sstate
1001	Dirty Rick	New York	NY
1002	Squirrels	Boston	MA

### PART

pno	pname	psize
999	Batteries	Large

### SUPPLIES

parent	child

### SUPPLY

qty	price
10	\$100
14	\$99

### SUPPLIED\_BY

parent	child

# Network Data Model

*Instance*

## SUPPLIER

sno	sname	scity	sstate
1001	Dirty Rick	New York	NY
1002	Squirrels	Boston	MA

## SUPPLIES

parent	child

## SUPPLY



qty	price
10	\$100
14	\$99

## PART

pno	pname	psize
999	Batteries	Large

## SUPPLIED\_BY

parent	child



# Network Data Model

## *Instance*

### SUPPLIER

sno	sname	scity	sstate
1001	Dirty Rick	New York	NY
1002	Squirrels	Boston	MA

### PART

pno	pname	psize
999	Batteries	Large

### SUPPLIES

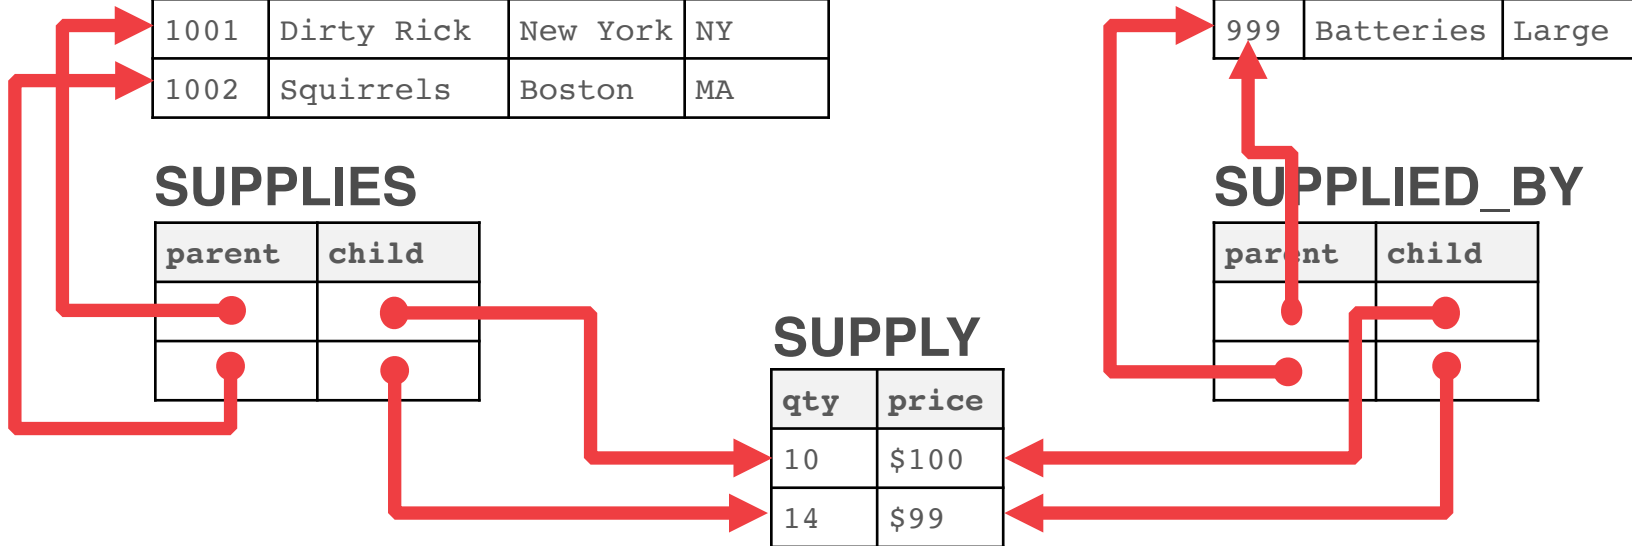
parent	child

### SUPPLIED\_BY

parent	child

### SUPPLY

qty	price
10	\$100
14	\$99



# 1960S – IBM IMS

---

## Information Management System

Early DBMS developed to keep track of purchase orders for Apollo moon mission.

- Hierarchical data model.
- Programmer-defined physical storage format.
- Tuple-at-a-time queries.



# Hierarchical Data Model

---

*Schema*

*Instance*

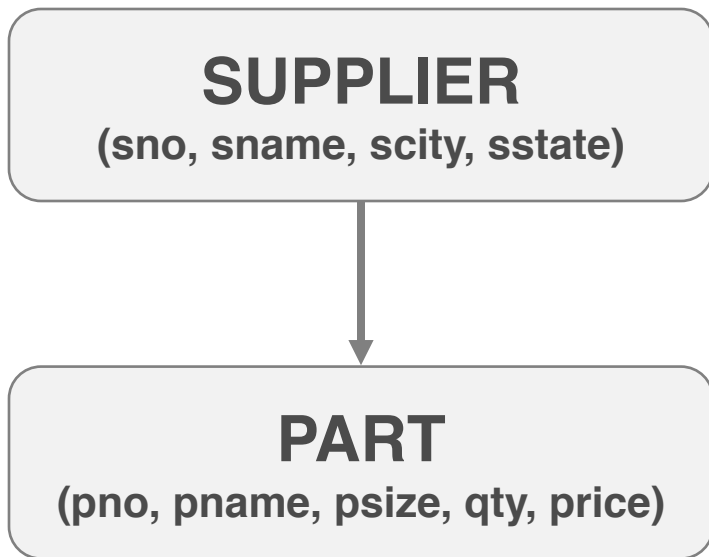
**SUPPLIER**  
(sno, sname, scity, sstate)



**PART**  
(pno, pname, psize, qty, price)

# Hierarchical Data Model

## *Schema*




## *Instance*

sno	sname	scity	sstate	parts
1001	Dirty Rick	New York	NY	
1002	Squirrels	Boston	MA	

# Hierarchical Data Model

## *Schema*



## *Instance*

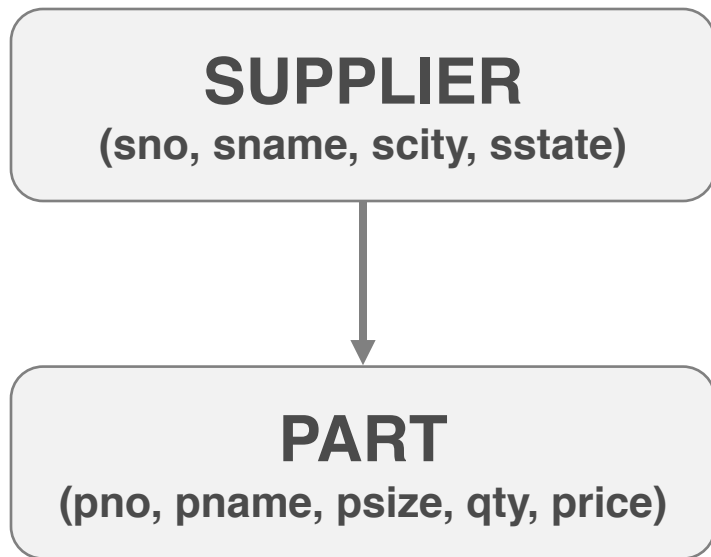
sno	sname	scity	sstate	parts
1001	Dirty Rick	New York	NY	999
1002	Squirrels	Boston	MA	

pno	pname	psize	qty	price
999	Batteries	Large	10	\$100

# Hierarchical Data Model

## *Schema*

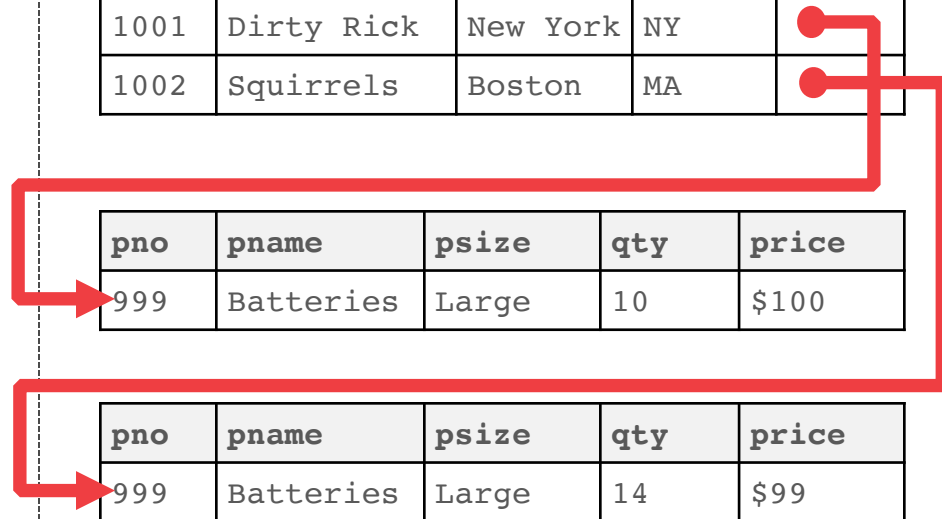


## *Instance*

sno	sname	scity	sstate	parts
1001	Dirty Rick	New York	NY	
1002	Squirrels	Boston	MA	

pno	pname	psize	qty	price
999	Batteries	Large	10	\$100

pno	pname	psize	qty	price
999	Batteries	Large	14	\$99






# 1970s – Relational Model

---

Ted Codd was a mathematician working at IBM Research. He saw developers spending their time rewriting IMS and CODASYL programs every time the database's schema or layout changed.

Database abstraction to avoid this maintenance:

- Store database in simple data structures.
- Access data through high-level language.
- Physical storage left up to implementation.



Codd

# Relational Model

## UNIVERSITY, REDUNDANCY AND CONSISTENCY OF RELATIONS STORED IN LARGE DATA BASES

E. F. Codd  
Research Division  
San Jose, California


**ABSTRACT:** The large, integrated data banks of the future will contain many relations of various degrees of redundancy. It will not be unusual for this set of stored relations to be redundant. Two types of redundancy are defined and discussed. The type may be employed to improve accessibility of certain kinds of information which happen to be in great demand. When either type of redundancy exists, those responsible for control of the data bank should know about it and have some means of detecting any "logical" inconsistencies in the total set of stored relations. Consistency checking might be helpful in tracking down anomalies (and possibly fraudulent) changes in the data bank contents.

83 0004 11113 August 30, 1969

**LIMITED DISTRIBUTION NOTICE:** - This report has been submitted for publication elsewhere and has been stored as a research report for early dissemination of its contents. As a courtesy to the intended publisher, it should not be widely distributed until after the date of version publication.

Copies may be obtained from IBM World's Library Service Center, Post Office Box 185, Yorktown Heights, New York 10598

mathematician working  
saw developers  
writing IMS and  
every time the  
layout changed.  
to avoid this



Codd

the data structures.  
high-level language.  
to implementation.



# Relational Data Model

## *Schema*



# Relational Data Model

## *Instance*

### SUPPLIER

sno	sname	scity	sstate
1001	Dirty Rick	New York	NY
1002	Squirrels	Boston	MA

### PART

pno	pname	psize
999	Batteries	Large

### SUPPLY

sno	pno	qty	price
1001	999	10	\$100
1002	999	14	\$99

# Relational Data Model

*Instance*

## SUPPLIER

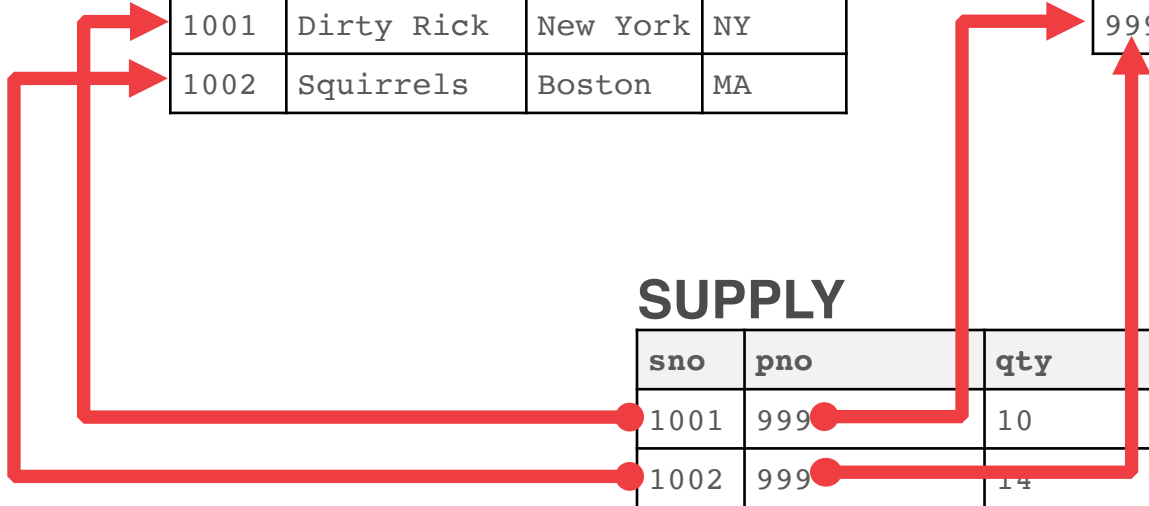
sno	sname	scity	sstate
1001	Dirty Rick	New York	NY
1002	Squirrels	Boston	MA

## PART

pno	pname	psize
999	Batteries	Large

## SUPPLY

sno	pno	qty	price
1001	999	10	\$100
1002	999	14	\$99



# 1970s – Relational Model

---

Early implementations of relational DBMS:

- **Peterlee Relational Test Vehicle** – IBM Research (UK)
- **System R** – IBM Research (San Jose)
- **INGRES** – U.C. Berkeley
- **Oracle** – Larry Ellison
- **Mimer** – Uppsala University



Gray



Stonebraker



Ellison

# 1980s – Relational Model

---

The relational model wins.

- IBM first releases SQL/DS in 1981.
- IBM then releases DB2 in 1983.
- "SEQUEL" becomes the standard (SQL) after supposedly Stonebraker refused to talk to the ANSI standards committee.

Many new "enterprise" DBMSs but Oracle wins marketplace.

Stonebraker creates Postgres as an "object-relational" DBMS.



ORACLE®

Informix®

TANDEM

SYBASE®

TERADATA

INGRES

InterBase®



# 1980s – Relational Model

The relational model wins.

- IBM first releases SQL/DS in 1981
- IBM then releases DB2 in 1982
- "SEQUEL" becomes the standard after supposedly Stonebraker to the ANSI standards committee

Many new "enterprise" DBMSs but Oracle wins marketplace

Stonebraker creates Postgres as an "object-relational" DBMS.

IBM DB2 ORACLE

But Ingres did not show up at the committee meetings because founder Mike Stonebraker detested the idea of having technology standards. Stonebraker was vocal about it. He thought they inhibited innovation and artificially restricted what got to the marketplace. Maybe so, but his hard-line position probably did not help his company. Don Deutsch, who served as chairman of the database committee, summed things up this way: "I tell you, QUEL was a much nicer language than SQL. No rational person would have chosen SQL instead of QUEL. . . . Ingres was stupid."

*So it was that the database committee began evaluating SQL*

# 1980s – Object-Oriented Databases

---

Avoid "relational-object impedance mismatch" by tightly coupling objects and database.

Few of these original DBMSs from the 1980s still exist today but many of the technologies exist in other forms (e.g., XML, JSON).

**VERSANT**   **ObjectStore**    **MarkLogic™**

# Object-Oriented Model

---

## *Application Code*

```
class Student {  
    int id;  
    String name;  
    String email;  
    String phone[];  
}
```

# Object-Oriented Model

---

## *Application Code*

```
class Student {  
  int id;  
  String name;  
  String email;  
  String phone[];  
}
```

## *Relational Schema*

**STUDENT**

(id, name, email)



**STUDENT\_PHONE**

(sid, phone)

# Object-Oriented Model

## *Application Code*

```
class Student {  
    int id;  
    String name;  
    String email;  
    String phone[];  
}
```

id	name	email
1001	M.O.P.	ante@up.com

sid	phone
1001	444-444-4444
1001	555-555-5555

## *Relational Schema*

**STUDENT**  
(id, name, email)



**STUDENT\_PHONE**  
(sid, phone)

# Object-Oriented Model

## *Application Code*

```
class Student {  
  int id;  
  String name;  
  String email;  
  String phone[];  
}
```

id	name	email
1001	M.O.P.	ante@up.com

sid	phone
1001	444-444-4444
1001	555-555-5555

## *Relational Schema*

**STUDENT**  
(id, name, email)



**STUDENT\_PHONE**  
(sid, phone)

# Object-Oriented Model

---

## *Application Code*

```
class Student {  
    int id;  
    String name;  
    String email;  
    String phone[];  
}
```

# Object-Oriented Model

---

## *Application Code*

```
class Student {  
    int id;  
    String name;  
    String email;  
    String phone[];  
}
```



### Student

```
{  
  "id": 1001,  
  "name": "M.O.P.",  
  "email": "ante@up.com",  
  "phone": [  
    "444-444-4444",  
    "555-555-5555"  
  ]  
}
```



# 1990s – Boring Years

---

No major advancements in database systems or application workloads.

- Microsoft forks Sybase and creates SQL Server.
- MySQL is written as a replacement for mSQL.
- Postgres gets SQL support.
- SQLite started in early 2000.

Some DBMSs introduced pre-computed [data cubes](#) for faster analytics.



## 2000s – Internet Boom

---

All the big players were heavyweight and expensive. Open-source DBMSs were missing important features.

Many companies wrote their own custom middleware to scale across many independent single-node DBMS instances.

# 2000s – Data Warehouses

---

Rise of the special purpose OLAP DBMSs.

- Distributed / Shared-Nothing
- Relational / SQL
- Usually closed-source

Significant performance benefits from using columnar data storage model.



# 2000s – MapReduce

---

Distributed programming and execution model for analyzing large data sets.

- First proposed by Google (**MapReduce**).
- Yahoo! created an open-source version (**Hadoop**).
- Data model decided by user-written functions.

People (eventually) realized this was a bad idea and grafted SQL on top of MapReduce. That was a bad idea too.



**MAPR-DB**

# 2000s – NoSQL

---

Focus on high-availability & high-scalability:

- Schema-less (i.e., "Schema Last")
- Non-relational data models (document, key/value, column-family)
- No ACID transactions
- Custom APIs instead of SQL
- Usually open-source



# 2010s – NewSQL

---

Provide same performance for OLTP as NoSQL without giving up ACID:

- Relational / SQL
- Distributed

Almost all of the first group of systems failed.

Second wave of "distributed SQL" systems are (potentially) doing better.



# 2010s – NewSQL

---

Provide same performance for OLTP as NoSQL without giving up ACID:

- Relational / SQL
- Distributed

Almost all of the first group of systems failed.

Second wave of "distributed SQL" systems are (potentially) doing better.



# 2010s – HTAP

---

## Hybrid Transactional-Analytical Processing

Execute fast OLTP like a NewSQL system  
while also executing complex OLAP queries  
like a data warehouse system.

- Distributed / Shared-Nothing
- Relational / SQL
- Mixed open/closed-source.





# 2010s – Stream Processing

---

Execute continuous queries on streams of tuples, extending semantics to include notion of windows.

Often used in combination with batch-oriented systems in a lambda architecture.



# 2010s – The Cloud

---

First database-as-a-service (DBaaS) offerings were "containerized" versions of existing DBMSs.

There are newer DBMSs that are designed from scratch explicitly for running in a cloud environment.



# 2010s – Shared-Disk

---

Instead of writing a custom storage manager, the DBMS leverages distributed storage.

- Scale execution layer independently of storage.
- Favors log-structured approaches.

This is what most people think of when they talk about a **data lake**.



# 2010s – Graphs

---

Systems for storing and querying graph data.

→ Similar to the network data model (CODASYL)

Their (supposed) advantage over other data models is to provide a graph-centric query API

→ [SQL:2023](#) is adding graph query syntax (SQL/PCG)

Latest [research](#) (2023) shows that relational DBMSs outperform graph DBMSs.



neo4j



MEM  
GRAPH



TigerGraph



NebulaGraph



JanusGraph



Dgraph




IndraDB



TerminusDB

graphbase.ai



APACHE  
GIRAPH

# 2010s – Time Series

---

Specialized systems that are designed to store time series / event data.

The design of these systems make deep assumptions about the distribution of data and workload query patterns.



# 2020s – Blockchains

---

Decentralized distributed log with incremental checksums ([Merkle Trees](#)).

→ Uses Byzantine Fault Tolerant (BFT) protocol to determine next entry to append to log.

Many blockchain use cases seem like they can be solved with a "traditional" OLTP DBMS and/or external policies (e.g., authentication).



**fluree**

**BIGCHAIN**DB



Condensation



**CovenantSQL**

# Current State of Affairs


---

The demarcation lines of DBMS categories will continue to blur over time as specialized systems expand the scope of their domains.

→ Every NoSQL DBMS (except for Redis) now supports SQL.

The relational model and declarative query languages promote better data engineering.

# DBMS Overview





# Distributed Query Execution

---

Executing OLAP queries in a distributed DBMS is roughly the same as on a single node.


→ Query plan is a DAG of physical operators.

For each operator, the DBMS considers where input is coming from and where to send output.


- Table Scans
- Joins
- Aggregations
- Sorting

# Distributed Query Execution

---



⋮




*Worker Nodes*

# Distributed Query Execution


---

*Persistent Data*




⋮

*Persistent Data*

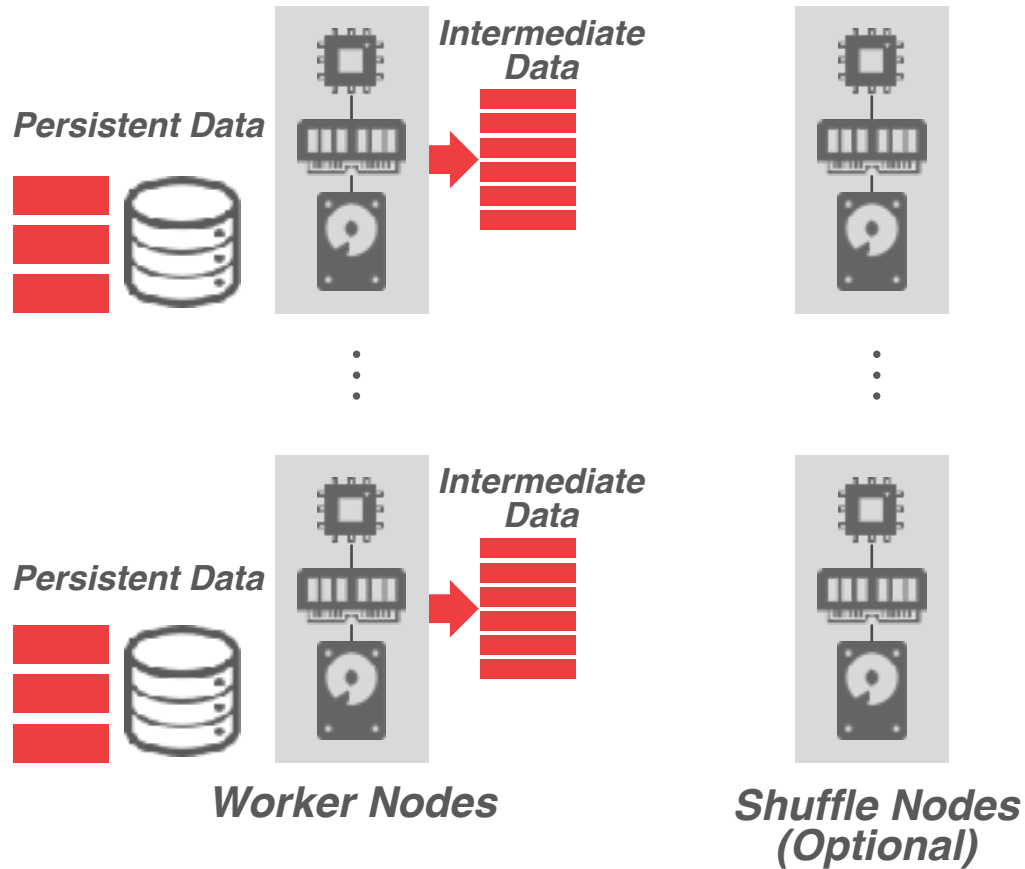


*Worker Nodes*

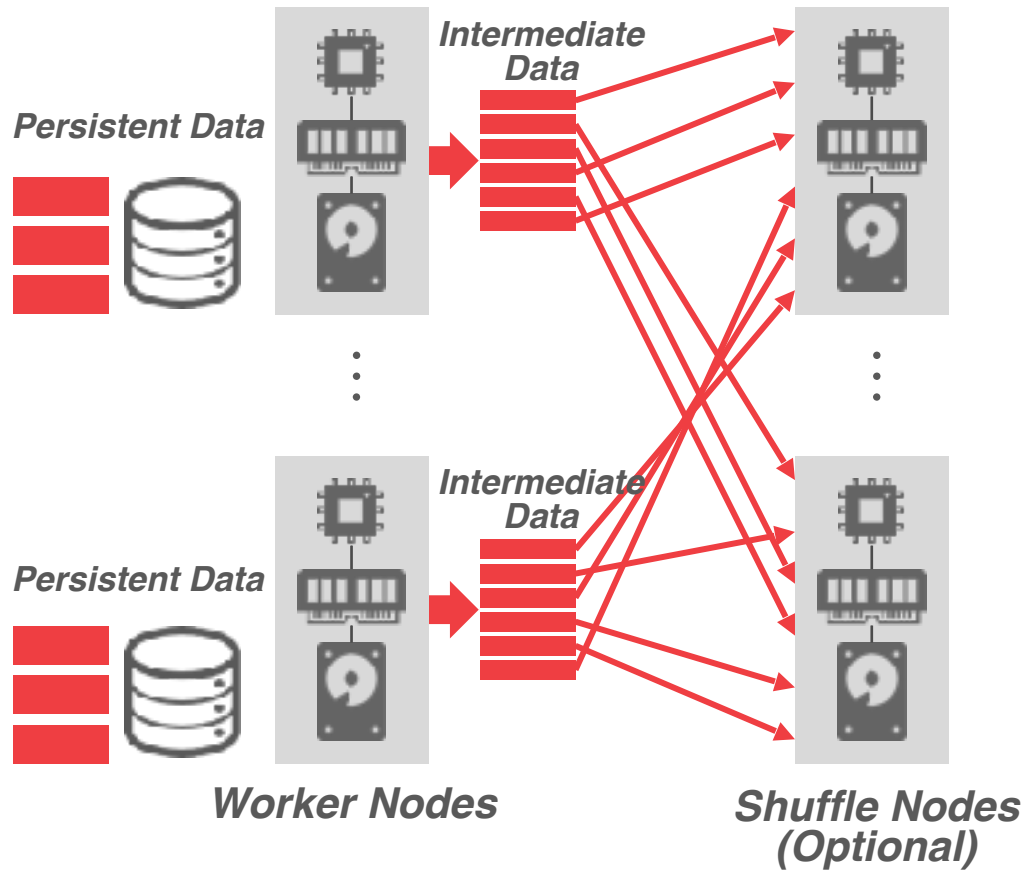
# Distributed Query Execution



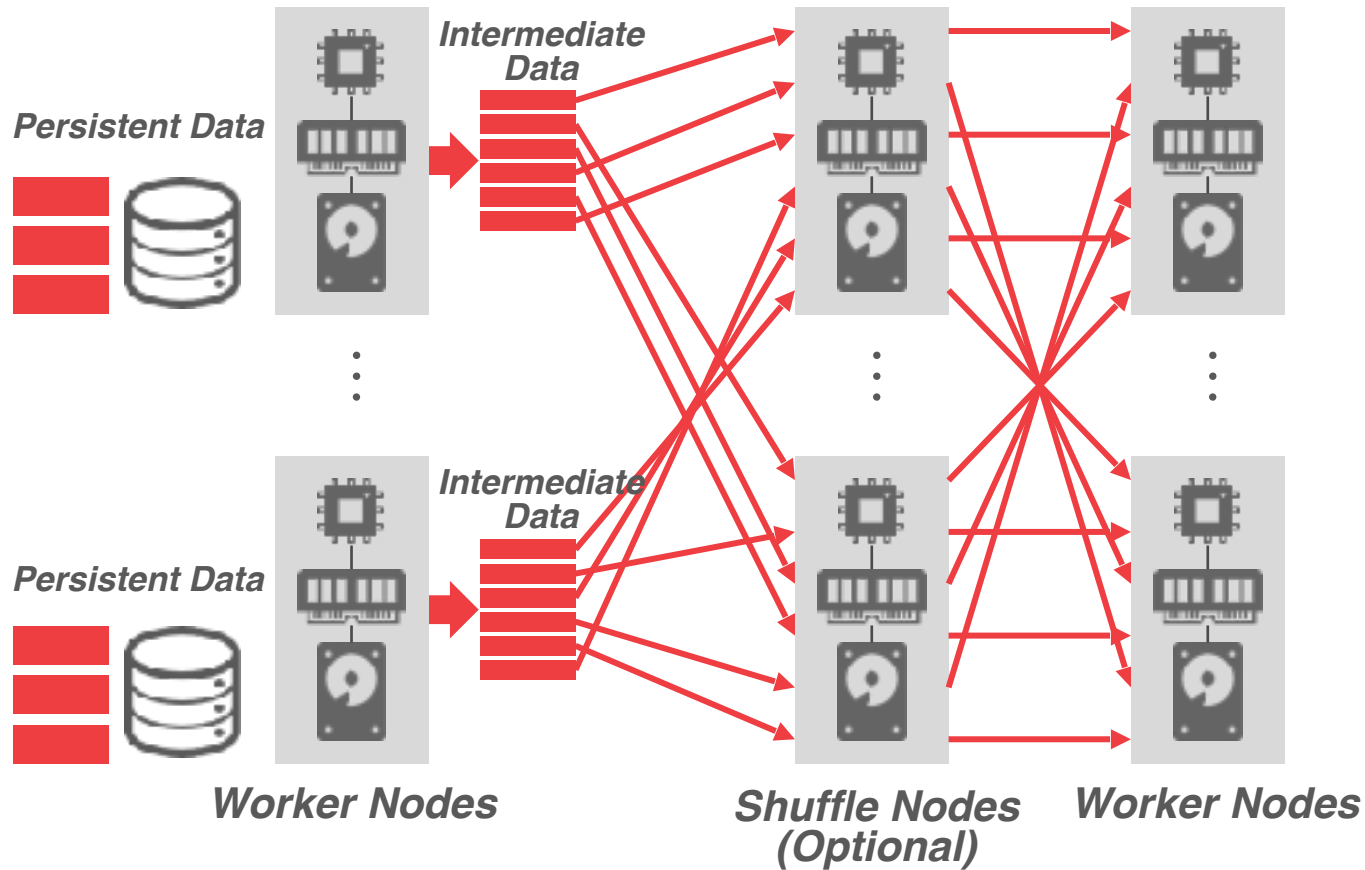
# Distributed Query Execution



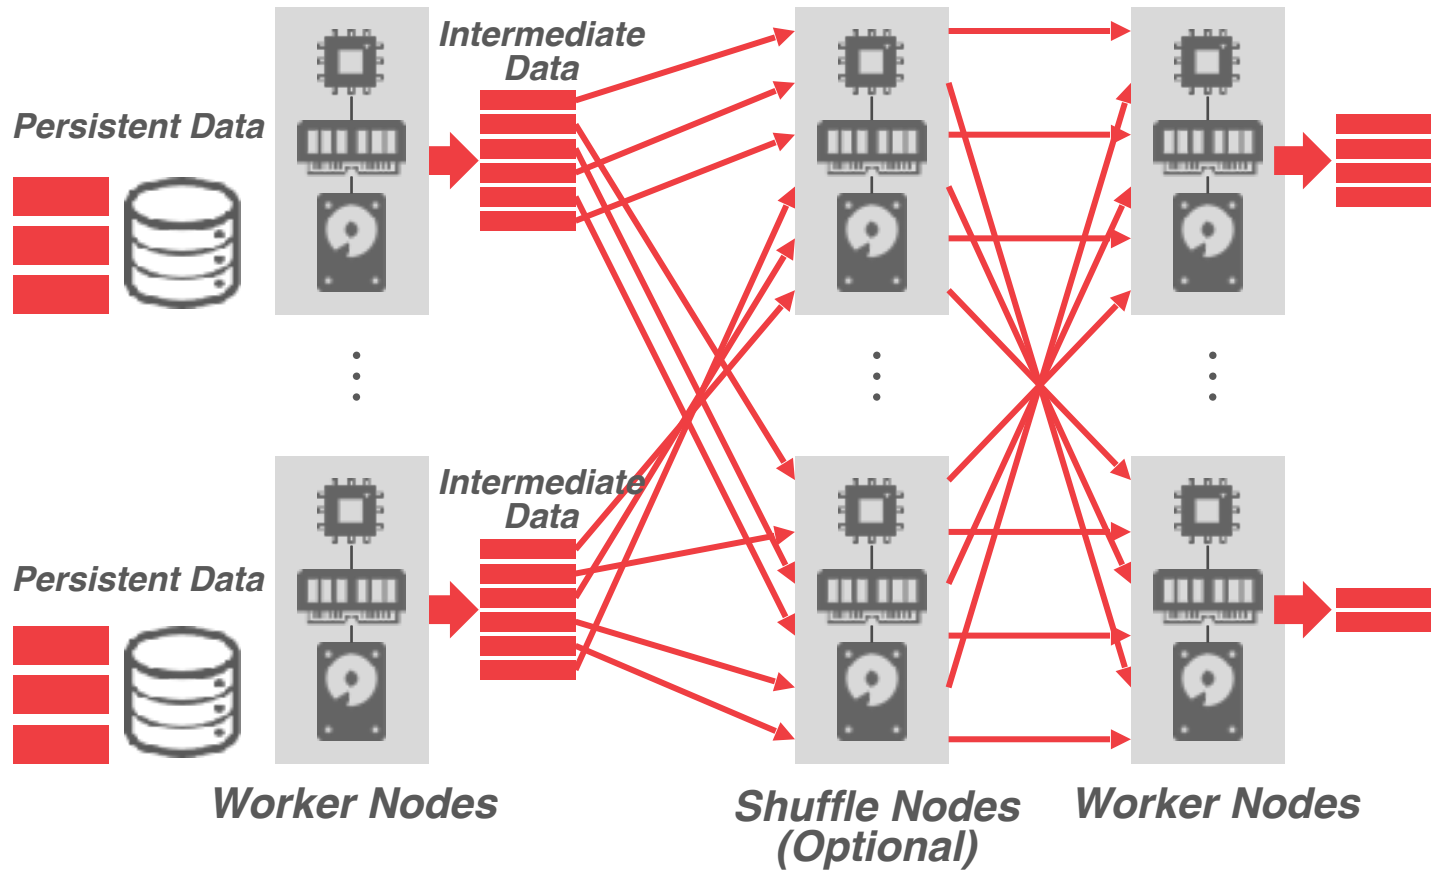
# Distributed Query Execution



# Distributed Query Execution

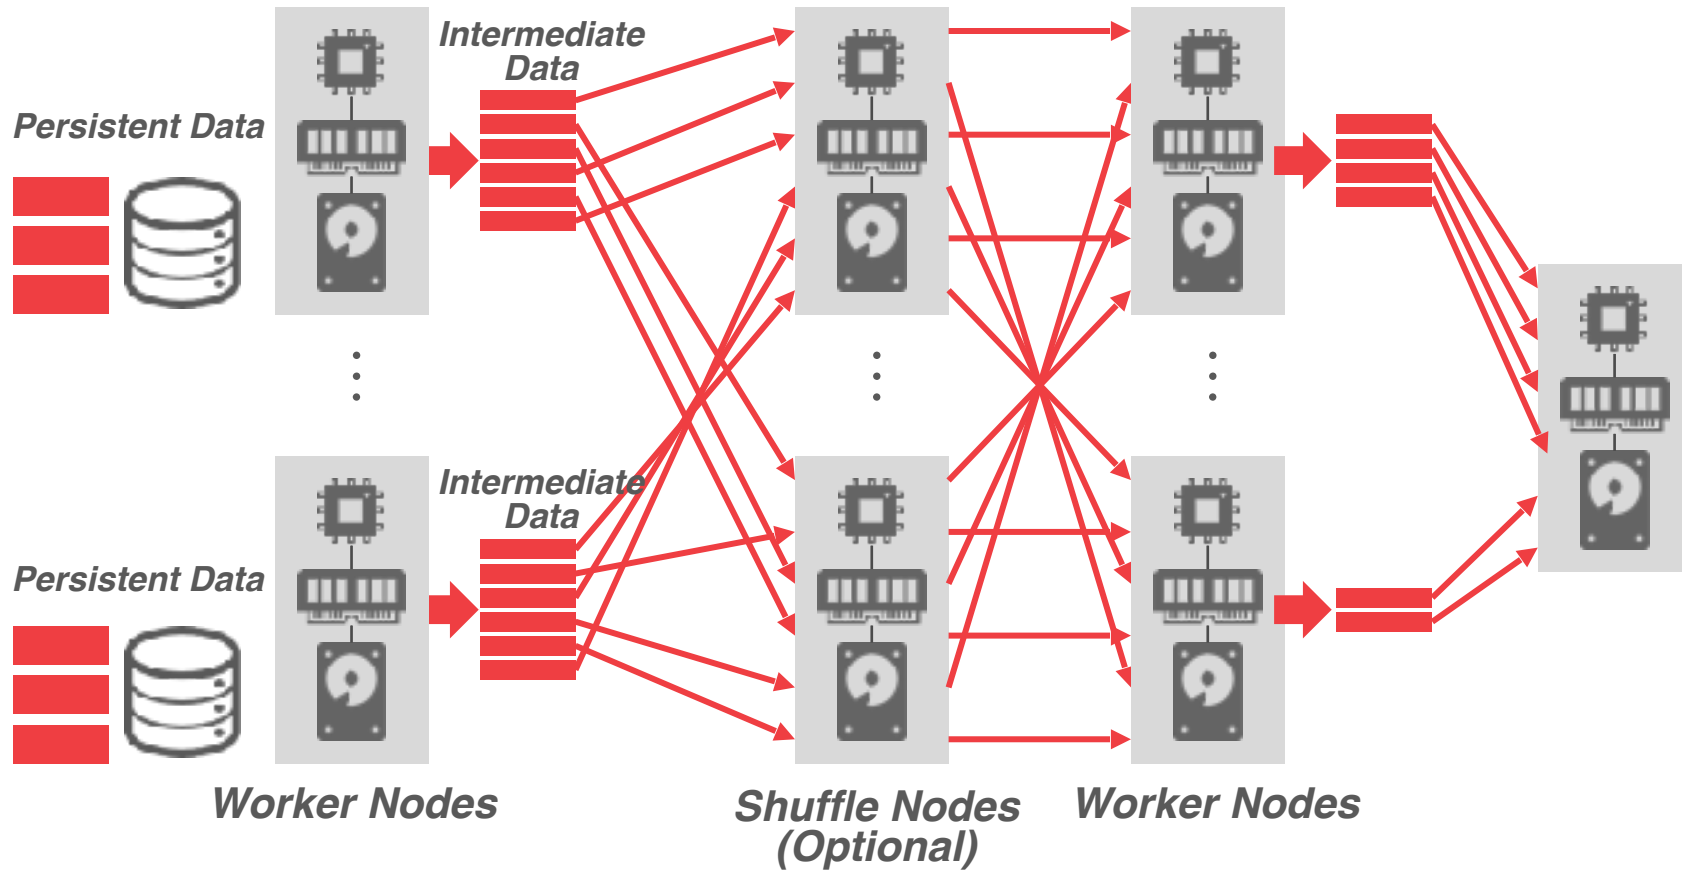


# Distributed Query Execution

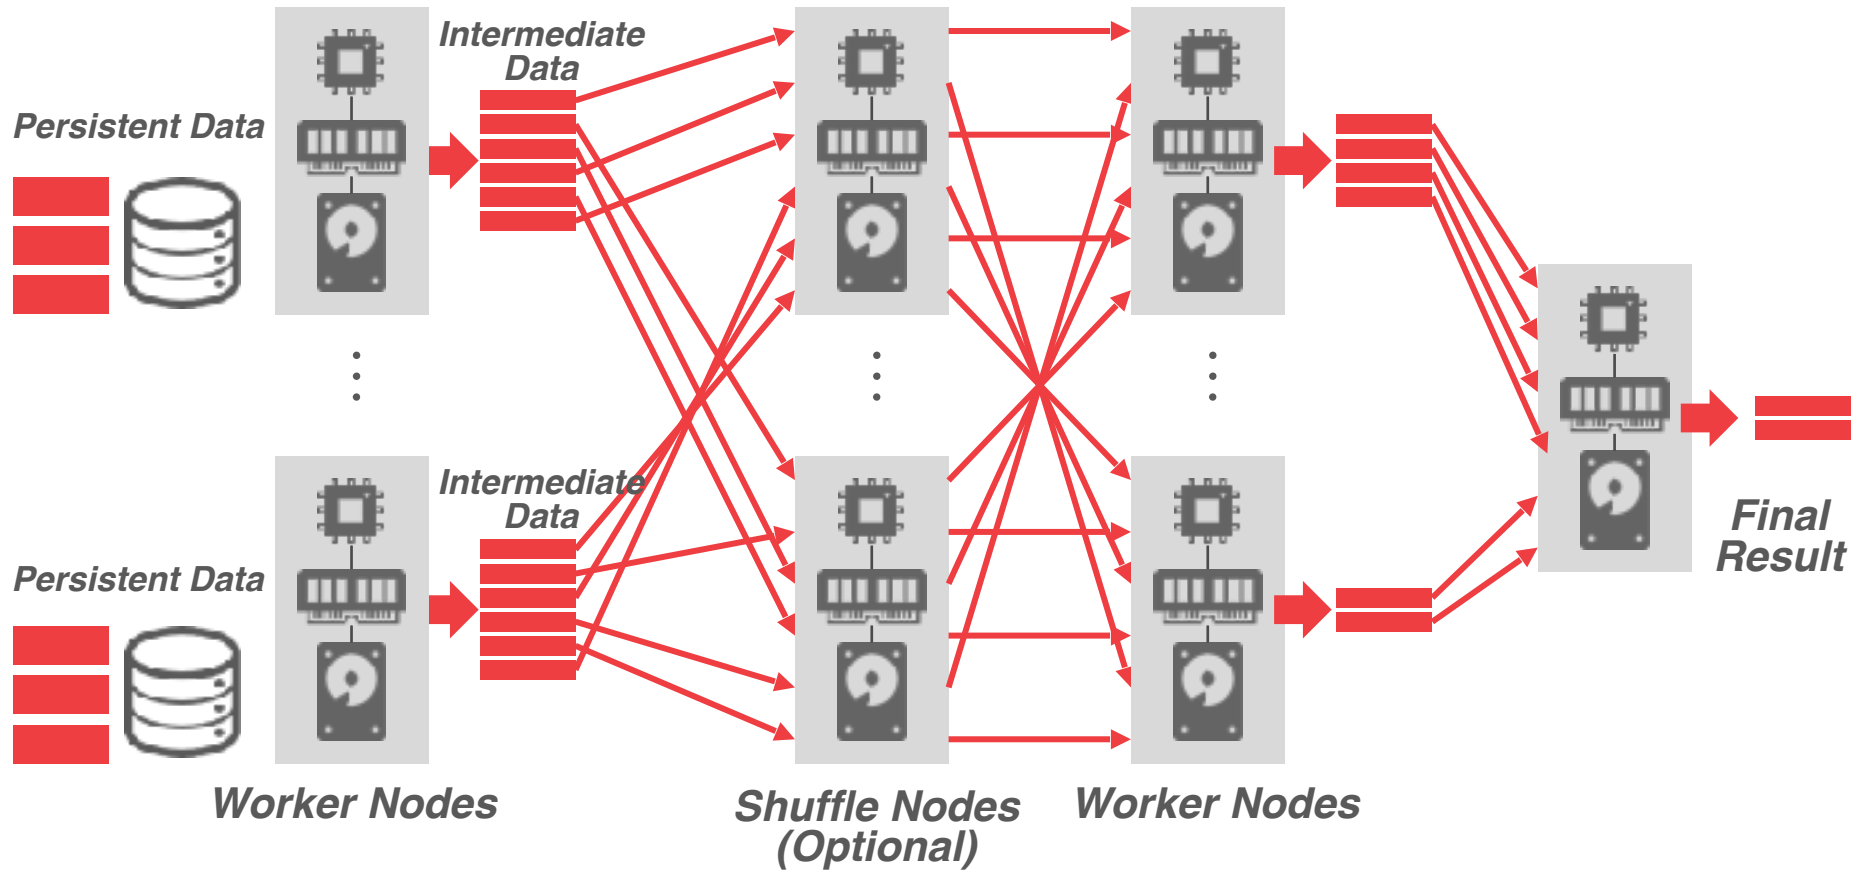




# Distributed Query Execution



# Distributed Query Execution



# Types of Data

---

## Persistent Data:

- The "source of record" for the database (e.g., tables).
- Modern systems assume that these data files are immutable but can support updates by rewriting them.

## Intermediate Data:

- Short-lived artifacts produced by query operators during execution and then consumed by other operators.
- The amount of intermediate data that a query generates has little correlation to amount of persistent data that it reads or the overall execution time.



# Distributed Architecture

---

A distributed DBMS's system architecture specifies the location of the database's persistent data files. This affects how nodes coordinate with each other and where they retrieve / store objects in the database.

Two approaches (not mutually exclusive):

- **Push Query to Data**
- **Pull Data to Query**

# Push vs. Pull

---

## **Approach #1: Push Query to Data**

- Send the query (or a portion of it) to the node that contains the data.
- Perform as much filtering and processing as possible where data resides before transmitting over network.

## **Approach #2: Pull Data to Query**

- Bring the data to the node that is executing a query that needs it for processing.
- This is necessary when there is no compute resources available where persistent data files are located.

## Approach 1

- Send the data to a location that contains the data.
- Perform the processing where the data is located.

## Approach 2

- Bring the data to a location that needs it for processing.
- This is necessary when there is no compute resources available where persistent data files are located.

### Filtering and retrieving data using Amazon S3 Select



[PDF](#) [RSS](#)

With Amazon S3 Select, you can use simple structured query language (SQL) statements to filter the contents of an Amazon S3 object and retrieve just the subset of data that you need. By using Amazon S3 Select to filter this data, you can reduce the amount of data that Amazon S3 transfers, which reduces the cost and latency to retrieve this data.

Amazon S3 Select works on objects stored in CSV, JSON, or Apache Parquet format. It also works with objects that are compressed with GZIP or BZIP2 (for CSV and JSON objects only), and server-side encrypted objects. You can specify the format of the results as either CSV or JSON, and you can determine how the records in the result are delimited.

You pass SQL expressions to Amazon S3 in the request. Amazon S3 Select supports a subset of SQL. For more information about the SQL elements that are supported by Amazon S3 Select, see [SQL reference for Amazon S3 Select](#).

You can perform SQL queries using AWS SDKs, the SELECT Object Content REST API, the AWS Command Line Interface (AWS CLI), or the Amazon S3 console. The Amazon S3 console limits the amount of data returned to 40 MB. To retrieve more data, use the AWS CLI or the API.

## Approach

### Filtering and retrieving data using Amazon S3 Select



PDF | RSS

With Amazon S3 Select, you can use SQL statements to filter the contents of an object at your need. By using Amazon S3 Select to filter this data, you can reduce the cost and latency to retrieve this data.



Feedback

### Query Blob Contents

Article • 07/20/2021 • 10 minutes to read • 3 contributors

The `query-blob-contents` API applies a simple Structured Query Language (SQL) statement on a blob's contents and returns only the queried subset of the data. You can also call `query-blob-contents` to query the contents of a version or snapshot.

### Request

The `query-blob-contents` request may be constructed as follows. HTTPS is recommended. Replace `myaccount` with the name of your storage account.

POST Method Request URI	HTTP Version
<code>https://myaccount.blob.core.windows.net/mycontainer/myblob?comp=query</code>	HTTP/1.0
<code>https://myaccount.blob.core.windows.net/mycontainer/myblob?comp=query&amp;snapshot=qustoid=</code>	HTTP/1.1
<code>https://myaccount.blob.core.windows.net/mycontainer/myblob?comp=query&amp;versionId=qustoid=</code>	

query language (SQL) statements to filter the contents of an object at your need. By using Amazon S3 Select to filter this data, you can reduce the cost and latency to retrieve this data.

Amazon S3 Select also works with objects that are in the Parquet or Apache Parquet format. It also works with objects that are server-side encrypted. You can specify the delimiter to determine how the records in the result are delimited.

Amazon S3 Select supports a subset of SQL. For more information about Amazon S3 Select, see [SQL reference for Amazon S3 Select](#).

Object Content REST API, the AWS Command Line Interface (CLI) limits the amount of data returned to 40 MB. To retrieve

to compute  
ent data files are

# Shared-Nothing


Each DBMS instance has its own CPU, memory, locally-attached disk.

→ Nodes only communicate with each other via network.

Database is partitioned into disjoint subsets across nodes.

→ Adding a new node requires physically moving data between nodes.

Since data is local, the DBMS can access it via POSIX API.





# Shared-Nothing


Each DBMS instance has its own CPU, memory, locally-attached disk.

→ Nodes only communicate with each other via network.

Database is partitioned into disjoint subsets across nodes.

→ Adding a new node requires physically moving data between nodes.

Since data is local, the DBMS can access it via POSIX API.

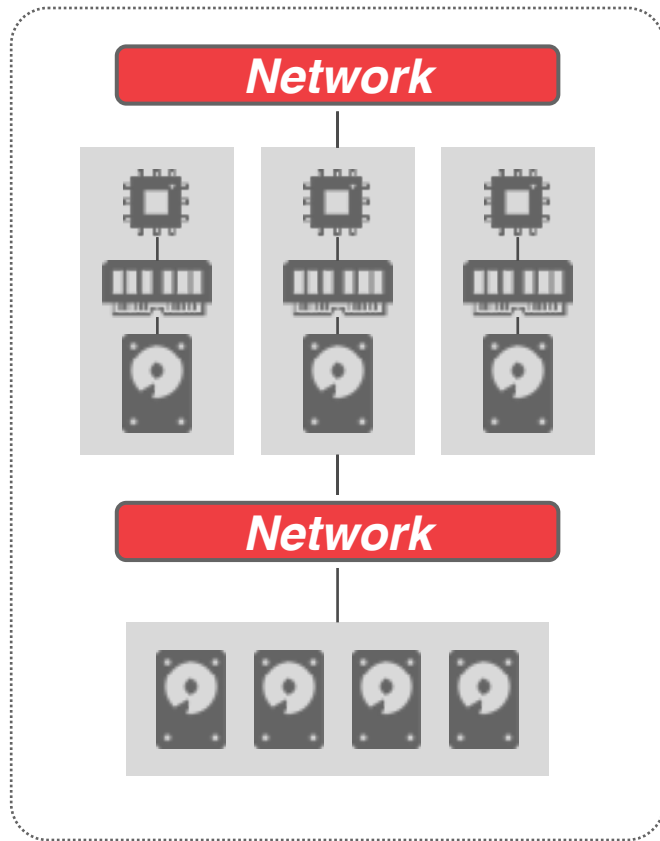


# Shared-Disk

Each node accesses a single logical disk via an interconnect, but also has its own private memory and ephemeral storage.

→ Must send messages between nodes to learn about their current state.

Instead of a POSIX API, the DBMS accesses disk via a userspace API.



# Shared-Disk

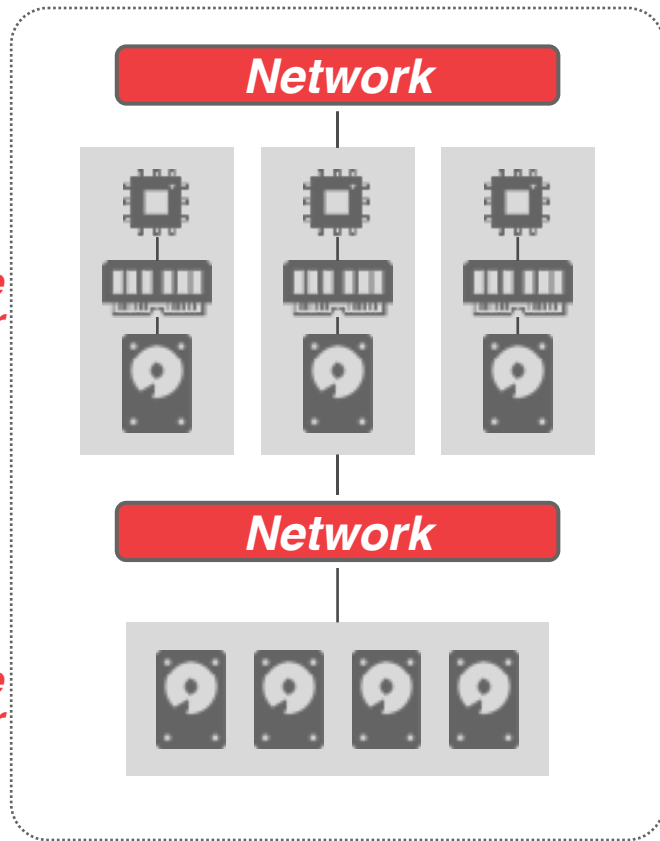
Each node accesses a single logical disk via an interconnect, but also has its own private memory and ephemeral storage.

→ Must send messages between nodes to learn about their current state.

*Compute Layer*

Instead of a POSIX API, the DBMS accesses disk via a userspace API.

*Storage Layer*



# System Architecture

---

## **Choice #1: Shared-Nothing:**

- Harder to scale capacity (data movement).
- Potentially better performance & efficiency.
- Apply filters where the data resides before transferring.

## **Choice #2: Shared-Disk:**

- Scale compute / storage layers independently.
- Easy to shut down idle compute layer resources.
- May need to pull uncached persistent data from storage layer to compute layer before filtering.

# System Architecture

---

## **Choice #1: Shared-Nothing:**

- Harder to scale capacity (data movement).
- Potentially better performance & efficiency.
- Apply filters where the data resides before transferring.

## **Choice #2: Shared-Disk:**

- Scale compute / storage layers independently.
- Easy to shut down idle compute layer resources.
- May need to pull uncached persistent data from storage layer to compute layer before filtering.

# Shared-Disk

---

Traditionally the storage layer in shared-disk DBMSs were dedicated on-prem NAS.

→ Example: Oracle Exadata

Cloud **object stores** are now the prevailing storage target for modern OLAP DBMSs because they are "infinitely" scalable.

→ Examples: Amazon S3, Azure Blob, Google Cloud Storage

# Object Stores

---

Partition the database's tables (persistent data) into large, immutable files stored in an object store.

- All attributes for a tuple are stored in the same file in a columnar layout (PAX).
- Header (or footer) contains meta-data about columnar offsets, compression schemes, indexes, and zone maps.

The DBMS retrieves a block's header to determine what byte ranges it needs to retrieve (if any).

Each cloud vendor provides their own proprietary API to access data (**PUT, GET, DELETE**).

- Some vendors support predicate pushdown (S3).

# Object Stores

Partiti  
into la

→ All a

colu

→ Hea

offs

The D

what


Each

API t

→ So

## Workers

- Separated compute / storage
  - One Worker pod per compute node
    - Executes portions of the query plan
- Custom network protocol over UDP
  - Data distribution between workers
  - Uses Intel DPDK
  - 50% higher throughput on AWS over TCP/IP
- Shard files cached in local NVMe SSD
- Shards persisted in object store
  - Custom AWS S3 access library
  - 3X better throughput than stock S3 lib





# Observation

---

Snowflake is a monolithic system comprised of components built entirely in-house.

Most of the non-academic DBMSs we will cover this semester will have a similar overall architecture.

But this means that multiple organizations are writing the same DBMS software...

# OLAP Commoditization

---

One recent trend in the last decade is the refactoring of OLAP engine sub-systems into standalone open-source components.

→ This is typically done by organizations not in the business of selling DBMS software.

## **Examples:**

- System Catalogs
- Query Optimizers
- File Format / Access Libraries
- Execution Engines

## Examples:

- System Catalogs
- Query Optimizers
- File Format / Access Libraries
- Execution Engines

*Permission is hereby granted by all the parties of this agreement, without prejudice, that the copies are free made or distributed for direct economic advantage. The TV-8 copyright owner understands the role of the publications and its direct impact and nature is given that support in the presentation of the New King's Book Project. Furthermore, it says otherwise, means random, renders a particular special permission from the Enforcement.*

*Proceedings of the 20th International Conference on Very Large Databases, Cairo, Egypt, 2000*

To begin our analyses, let us put together a few important observations: on how Internet systems are viewed by economic vendors, and the research community.

*Pharmaceutical 3* formulates drug products to ensure compatibility. Database systems offer more and more features, making it difficult to find and use studies.

# System Catalogs

---

A DBMS tracks a database's schema (e.g., table, columns) and data files in its catalog.

- If the DBMS is on the data ingestion path, then it can maintain the catalog incrementally.
- If an external process adds data files, then it also needs to update the catalog so that the DBMS is aware of them.

Notable implementations:

- [HCatalog](#)
- [Google Data Catalog](#)
- [Amazon Glue Data Catalog](#)

# Query Optimizers

---

Extendible search engine framework for heuristic- and cost-based query optimization.

- DBMS provides transformation rules and cost estimates.
- Framework returns either a logical or physical query plan.

This is the hardest part to build in any DBMS.

Notable implementations:

- [Greenplum Orca](#)
- [Apache Calcite](#)



# File Formats

---

Most DBMSs use a proprietary on-disk binary file format for their databases. The only way to share data between systems is to convert data into a common text-based format

→ Examples: CSV, JSON, XML

There are open-source binary file formats that make it easier to access data across systems and libraries for extracting data from files.

→ Libraries provide an iterator interface to retrieve (batched) columns from files.

# Storage Formats

---

## Apache Parquet (2013)

→ Compressed columnar storage from Cloudera/Twitter

## Apache ORC (2013)

→ Compressed columnar storage from Apache Hive.

## Apache CarbonData (2013)

→ Compressed columnar storage with indexes from Huawei.

## Apache Iceberg (2017)

→ Flexible data format that supports schema evolution from Netflix.

## HDF5 (1998)

→ Multi-dimensional arrays for scientific workloads.

## Apache Arrow (2016)

→ In-memory compressed columnar storage from Pandas / Dremio.

# Execution Engines

---

Standalone libraries for executing vectorized query operators on columnar data.

- Input is a DAG of physical operators.
- Require external scheduling and orchestration.

Notable implementations:

- [Velox](#)
- [DataFusion](#)
- [Intel OAP](#)



# Conclusion

---

Today was about understanding the high-level history and concept of modern OLAP DBMSs.

→ Fundamentally, they are not very different from previous distributed / parallel DBMSs except for their use of cloud-based object stores.

Our focus for the rest of the quarter will be on state-of-the-art implementations of the various system components.

# Next Class

---

Data Storage