

# 04: Data Skipping

Andrew Crotty // CS497 // Fall 2023

# Observation

---

- OLTP DBMSs use indexes to find individual tuples without performing sequential scans.
- Tree-based indexes (B+Trees) are meant for queries with low selectivity predicates.
  - Also need to accommodate incremental updates.

But OLAP queries usually don't need to find individual tuples and data files are read-only.

How can we speed up sequential scans?

# Data Skipping

---

## Approach #1: Approximate Queries (Lossy)

- Execute queries on a sampled subset of the entire table to produce approximate results.
- Examples: [BlinkDB](#), [Redshift](#), [ComputeDB](#), [XDB](#), [Oracle](#), [Snowflake](#), [Google BigQuery](#), [DataBricks](#)

## Approach #2: Data Pruning (Lossless)

- Use auxiliary data structures for evaluating predicates to quickly identify portions of a table that the DBMS can skip instead of examining tuples individually.
- DBMS must consider trade-offs between **scope** vs. **filter efficacy**, **manual** vs. **automatic**.

# Data Skipping

---

## Approach #1: Approximate Queries (Lossy)

- Execute queries on a sampled subset of the entire table to produce approximate results.
- Examples: [BlinkDB](#), [Redshift](#), [ComputeDB](#), [XDB](#), [Oracle](#), [Snowflake](#), [Google BigQuery](#), [DataBricks](#)

## Approach #2: Data Pruning (Lossless)

- Use auxiliary data structures for evaluating predicates to quickly identify portions of a table that the DBMS can skip instead of examining tuples individually.
- DBMS must consider trade-offs between **scope** vs. **filter efficacy**, **manual** vs. **automatic**.

# Data Considerations

---

## **Predicate Selectivity**

→ How many tuples will satisfy a query's predicates.

## **Skewness**

→ Whether a column has all unique values or contains many repeated values.

## **Clustering / Sorting**

→ Whether the table is pre-sorted on the attributes accessed in a query's predicates.

# Today's Agenda

---

Zone Maps

Bitmap Indexes

Bit-Slicing

Bit-Weaving

Column Imprints

Column Sketches

# Zone Maps

Pre-computed aggregates for a block of tuples. DBMS checks the zone map first to decide whether to access the block.

- Originally ***Small Materialized Aggregates*** (SMA)
- DBMS automatically maintains this meta-data.

## *Original Data*

val
100
200
300
400
400



# Zone Maps

Pre-computed aggregates for a block of tuples. DBMS checks the zone map first to decide whether to access the block.

- Originally **Small Materialized Aggregates** (SMA)
- DBMS automatically maintains this meta-data.

*Original Data*

val
100
200
300
400
400



*Zone Map*

type	val
MIN	100
MAX	400
AVG	280
SUM	1400
COUNT	5





# Zone Maps

Pre-computed aggregates for a block of tuples. DBMS checks the zone map first to decide whether to access the block.

- Originally ***Small Materialized Aggregates*** (SMA)
- DBMS automatically maintains this meta-data.

**SELECT \* FROM table  
WHERE val > 600**

***Original Data***

val
100
200
300
400
400



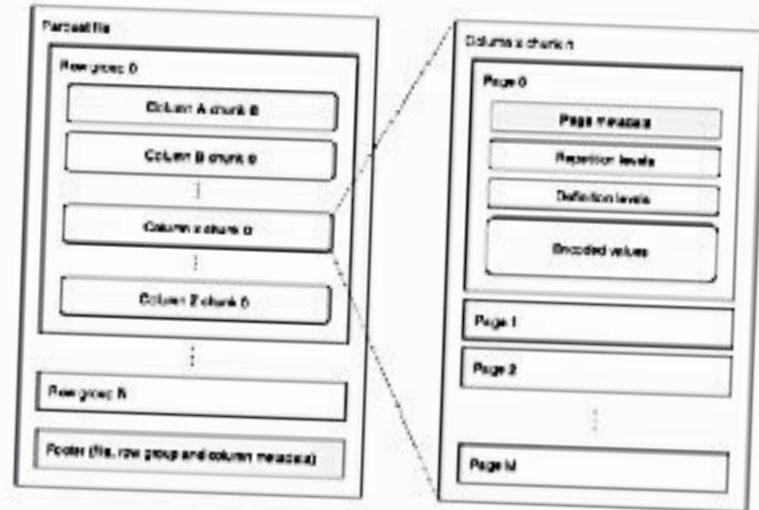
***Zone Map***

type	val
MIN	100
MAX	400
AVG	280
SUM	1400
COUNT	5




# Parquet: data organization

- Data organization
  - Row-groups (default 128MB)
  - Column chunks
  - Pages (default 1MB)
    - Metadata
      - Min
      - Max
      - Count
    - Rep/def levels
    - Encoded values



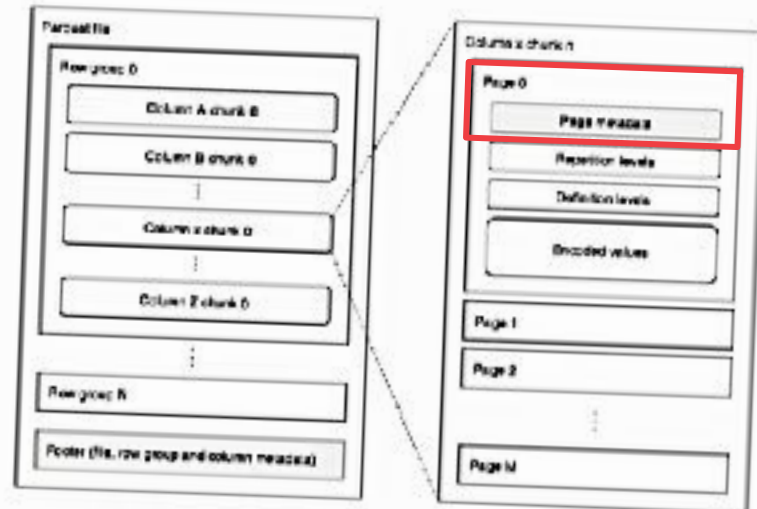
SELECT  
WHERE



SMALL MA  
LIGHT WEIGHT INDEX STRUCTURE FOR  
DATA WAREHOUSING  
VLDB 1998

# Parquet: data organization

- Data organization
  - Row-groups (default 128MB)
  - Column chunks
  - Pages (default 1MB)
    - Metadata
      - Min
      - Max
      - Count
    - Rep/def levels
    - Encoded values



SELECT  
WHERE



SMALL MA  
LIGHT WEIGHT INDEX STRUCTURE FOR  
DATA WAREHOUSING  
VLDB 1998

# Observation

---

Trade-off between scope vs. filter efficacy.

- If the scope is too large, then the zone maps will be useless.
- If the scope is too small, then the DBMS will spend too much time checking zone maps.

Zone Maps are only useful when the target attribute's position and values are correlated.

# Bitmap Indexes


---

Store a separate Bitmap for each unique value for an attribute where an offset in the vector corresponds to a tuple.

→ The  $i^{th}$  position in the Bitmap corresponds to the  $i^{th}$  tuple in the table.

Typically segmented into chunks to avoid allocating large blocks of contiguous memory.

→ Example: One per row group in PAX.



# Bitmap Indexes

---

## *Original Data*

id	lit?
1	Y
2	Y
3	Y
4	N
6	Y
7	N
8	Y
9	Y

# Bitmap Indexes

---

## *Original Data*

id	lit?
1	Y
2	Y
3	Y
4	N
6	Y
7	N
8	Y
9	Y

# Bitmap Indexes

*Original Data*

id	lit?
1	Y
2	Y
3	Y
4	N
6	Y
7	N
8	Y
9	Y



*Compressed Data*

id	lit?	
	Y	N
1	1	0
2	1	0
3	1	0
4	0	1
6	1	0
7	0	1
8	1	0
9	1	0



# Bitmap Indexes

*Original Data*

id	lit?
1	Y
2	Y
3	Y
4	N
6	Y
7	N
8	Y
9	Y



*Compressed Data*

id	lit?	
	Y	N
1	1	0
2	1	0
3	1	0
4	0	1
6	1	0
7	0	1
8	1	0
9	1	0

# Bitmap Indexes: Example

---

```
CREATE TABLE customer_dim  
(  
  id INT PRIMARY KEY,  
  name VARCHAR(32),  
  email VARCHAR(64),  
  address VARCHAR(64),  
  zip_code INT  
);
```

# Bitmap Indexes: Example

---

```
CREATE TABLE customer_dim  
(  
  id INT PRIMARY KEY,  
  name VARCHAR(32),  
  email VARCHAR(64),  
  address VARCHAR(64),  
  zip_code INT  
);
```

```
SELECT id, email FROM customer_dim  
WHERE zip_code IN  
(15216,15217,15218);
```

# Bitmap Indexes: Example

---

```
CREATE TABLE customer_dim  
(  
  id INT PRIMARY KEY,  
  name VARCHAR(32),  
  email VARCHAR(64),  
  address VARCHAR(64),  
  zip_code INT  
);
```

```
SELECT id, email FROM customer_dim  
WHERE zip_code IN  
(15216,15217,15218);
```

# Bitmap Indexes: Example

---

```
CREATE TABLE customer_dim  
(  
  id INT PRIMARY KEY,  
  name VARCHAR(32),  
  email VARCHAR(64),  
  address VARCHAR(64),  
  zip_code INT  
);
```

```
SELECT id, email FROM customer_dim  
WHERE zip_code IN  
(15216,15217,15218);
```

# Bitmap Indexes: Example

---

```
CREATE TABLE customer_dim  
(  
  id INT PRIMARY KEY,  
  name VARCHAR(32),  
  email VARCHAR(64),  
  address VARCHAR(64),  
  zip_code INT  
);
```

```
SELECT id, email FROM customer_dim  
WHERE zip_code IN  
(15216,15217,15218);
```

Take the intersection of three bitmaps to find matching tuples.

# Bitmap Indexes: Example

```
CREATE TABLE customer_dim  
(  
  id INT PRIMARY KEY,  
  name VARCHAR(32),  
  email VARCHAR(64),  
  address VARCHAR(64),  
  zip_code INT  
);
```

```
SELECT id, email FROM customer_dim  
WHERE zip_code IN  
(15216,15217,15218);
```

Take the intersection of three bitmaps to find matching tuples.

Assume we have 10 million tuples.

43,000 zip codes in the US.

→  $10000000 \times 43000 = 53.75\text{GB}$

# Bitmap Indexes: Example

```
CREATE TABLE customer_dim  
(  
  id INT PRIMARY KEY,  
  name VARCHAR(32),  
  email VARCHAR(64),  
  address VARCHAR(64),  
  zip_code INT  
);
```

```
SELECT id, email FROM customer_dim  
WHERE zip_code IN  
(15216,15217,15218);
```

Take the intersection of three bitmaps to find matching tuples.

Assume we have 10 million tuples.

43,000 zip codes in the US.

→  $10000000 \times 43000 = 53.75\text{GB}$



# Bitmap Indexes: Example

```
CREATE TABLE customer_dim  
(  
  id INT PRIMARY KEY,  
  name VARCHAR(32),  
  email VARCHAR(64),  
  address VARCHAR(64),  
  zip_code INT  
);
```

```
SELECT id, email FROM customer_dim  
WHERE zip_code IN  
(15216,15217,15218);
```

Take the intersection of three bitmaps to find matching tuples.

Assume we have 10 million tuples.

43,000 zip codes in the US.

→  $10000000 \times 43000 = 53.75\text{GB}$

This is **wasteful** because most entries in the bitmaps will be zeros.

→ Original:  $10000000 \times 32\text{-bit} = 40\text{MB}$

# Bitmap Indexes: Design Choices

---

## Encoding Scheme

→ How to represent and organize data in a Bitmap.

## Compression

→ How to reduce the size of sparse Bitmaps.

# Bitmap Indexes: Design Choices

---

## Encoding Scheme

→ How to represent and organize data in a Bitmap.

## Compression

→ How to reduce the size of sparse Bitmaps.

# Bitmap Indexes: Encoding

---

## **Approach #1: Equality Encoding**

→ Basic scheme with one Bitmap per unique value.

## **Approach #2: Range Encoding**

→ Use one Bitmap per interval instead of one per value.

## **Approach #3: Hierarchical Encoding**


→ Use a tree to identify empty key ranges.

## **Approach #4: Bit-sliced Encoding**

→ Use a Bitmap per bit location across all values.


# Hierarchical Encoding

*Offsets: 1, 3, 9, 12, 13, 14, 38, 40*




# Hierarchical Encoding

*Offsets: 1, 3, 9, 12, 13, 14, 38, 40*



# Hierarchical Encoding

*Offsets: 1, 3, 9, 12, 13, 14, 38, 40*




# Hierarchical Encoding

**Offsets:** 1, 3, 9, 12, 13, 14, 38, 40

**Original:** 8 bytes

**Encoded:** 4 bytes






# Bit-Sliced Encoding

---

## *Original Data*

id	zipcode
1	21042
2	15217
3	02903
4	90220
6	14623
7	53703

## *Bit-Slices*




# Bit-Sliced Encoding

---

## *Original Data*

id	zipcode
1	21042
2	15217
3	02903
4	90220
6	14623
7	53703

## *Bit-Slices*



# Bit-Sliced Encoding

## Original Data

id	zipcode
1	21042
2	15217
3	02903
4	90220
6	14623
7	53703



## Bit-Slices

$\text{bin}(21042) \rightarrow 00101001000110010$

# Bit-Sliced Encoding

## Original Data

id	zipcode
1	21042
2	15217
3	02903
4	90220
6	14623
7	53703



## Bit-Slices

null	16	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
------	----	----	----	----	----	----	----	---	---	---	---	---	---	---	---	---	---

$\text{bin}(21042) \rightarrow 00101001000110010$

# Bit-Sliced Encoding

## Original Data

id	zipcode
1	21042
2	15217
3	02903
4	90220
6	14623
7	53703



## Bit-Slices

16	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
0																

$\text{bin}(21042) \rightarrow 00101001000110010$

# Bit-Sliced Encoding

## Original Data

id	zipcode
1	21042
2	15217
3	02903
4	90220
6	14623
7	53703



## Bit-Slices

null	16	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
0																	



bin(21042) → 00101001000110010

# Bit-Sliced Encoding

## Original Data

id	zipcode
1	21042
2	15217
3	02903
4	90220
6	14623
7	53703



## Bit-Slices

	16	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
null	0	0	1	0	1	0	0	1	0	0	0	1	1	0	0	1	0



$\text{bin}(21042) \rightarrow 00101001000110010$

# Bit-Sliced Encoding

## Original Data

id	zipcode
1	21042
2	15217
3	02903
4	90220
6	14623
7	53703



## Bit-Slices

	16	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
null	0	0	1	0	1	0	0	1	0	0	0	1	1	0	0	1	0



# Bit-Sliced Encoding

## Original Data

id	zipcode
1	21042
2	15217
3	02903
4	90220
6	14623
7	53703



## Bit-Slices

null	16	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
0	0	0	1	0	1	0	0	1	0	0	0	1	1	0	0	1	0
0	0	0	0	1	1	1	0	1	1	0	1	1	1	0	0	0	1
0	0	0	0	0	0	1	0	1	1	0	1	0	1	0	1	1	1
0	1	0	1	1	0	0	0	0	0	0	1	1	0	1	1	0	0
0	0	0	0	1	1	1	0	0	1	0	0	0	1	1	1	1	1
0	0	1	1	0	1	0	0	0	1	1	1	0	0	0	1	1	1

# Bit-Sliced Encoding

## Original Data

id	zipcode
1	21042
2	15217
3	02903
4	90220
6	14623
7	53703



## Bit-Slices

null	16	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
0	0	0	1	0	1	0	0	1	0	0	0	1	1	0	0	1	0
0	0	0	0	1	1	1	0	1	1	0	1	1	1	0	0	0	1
0	0	0	0	0	0	1	0	1	1	0	1	0	1	0	1	1	1
0	1	0	1	1	0	0	0	0	0	0	1	1	0	1	1	0	0
0	0	0	0	1	1	1	0	0	1	0	0	0	1	1	1	1	1
0	0	1	1	0	1	0	0	0	1	1	1	0	0	0	1	1	1

```
SELECT * FROM
customer_dim
WHERE zipcode < 15217
```

# Bit-Sliced Encoding

## Original Data

id	zipcode
1	21042
2	15217
3	02903
4	90220
6	14623
7	53703



## Bit-Slices

null	16	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
0	0	0	1	0	1	0	0	1	0	0	0	1	1	0	0	1	0
0	0	0	0	1	1	1	0	1	1	0	1	1	1	0	0	0	1
0	0	0	0	0	0	1	0	1	1	0	1	0	1	0	1	1	1
0	1	0	1	1	0	0	0	0	0	0	1	1	0	1	1	0	0
0	0	0	0	1	1	1	0	0	1	0	0	0	1	1	1	1	1
0	0	1	1	0	1	0	0	0	1	1	1	0	0	0	1	1	1

```
SELECT * FROM
customer_dim
WHERE zipcode < 15217
```

Walk each slice and construct a result bitmap.

# Bit-Sliced Encoding

## Original Data

id	zipcode
1	21042
2	15217
3	02903
4	90220
6	14623
7	53703



## Bit-Slices

null	16	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
0	0	0	1	0	1	0	0	1	0	0	0	1	1	0	0	1	0
0	0	0	0	1	1	1	0	1	1	0	1	1	1	0	0	0	1
0	0	0	0	0	0	1	0	1	1	0	1	0	1	0	1	1	1
0	1	0	1	1	0	0	0	0	0	0	1	1	0	1	1	0	0
0	0	0	0	1	1	1	0	0	1	0	0	0	1	1	1	1	1
0	0	1	1	0	1	0	0	0	1	1	1	0	0	0	1	1	1

```
SELECT * FROM
customer_dim
WHERE zipcode < 15217
```

Walk each slice and construct a result bitmap.

# Bit-Sliced Encoding

## Original Data

id	zipcode
1	21042
2	15217
3	02903
4	90220
6	14623
7	53703



## Bit-Slices

null	16	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
0	0	0	1	0	1	0	0	1	0	0	0	1	1	0	0	1	0
0	0	0	0	1	1	1	0	1	1	0	1	1	1	0	0	0	1
0	0	0	0	0	0	1	0	1	1	0	1	0	1	0	1	1	1
0	1	0	1	1	0	0	0	0	0	0	1	1	0	1	1	0	0
0	0	0	0	1	1	1	0	0	1	0	0	0	1	1	1	1	1
0	0	1	1	0	1	0	0	0	1	1	1	0	0	0	1	1	1

```
SELECT * FROM
customer_dim
WHERE zipcode < 15217
```

Walk each slice and construct a result bitmap.  
Skip entries that have 1 in first 3 slices (16, 15, 14)

# Bit-Sliced Encoding

---

Bit-slices can also be used for efficient aggregate computations.

Example: **SUM(attr)** using Hamming Weight

- First, count the number of **1**s in **slice**<sub>17</sub> and multiply the count by  $2^{17}$
- Then, count the number of **1**s in **slice**<sub>16</sub> and multiply the count by  $2^{16}$
- Repeat for the rest of slices...

Intel added **POPCNT** SIMD instruction in 2008.

# Bitweaving

---

Alternative storage layout for columnar databases that is designed for efficient predicate evaluation on compressed data using SIMD.

- Order-preserving dictionary encoding.
- Bit-level parallelization.
- Only require common instructions (no scatter/gather)

Implemented in Wisconsin's [QuickStep](#) engine.

- Became an [Apache Incubator](#) project in 2016 but then died in 2018.



# Bitweaving

Alternative storage layout for columnar databases that is designed for efficient predicate evaluation on compressed data using SIMD.

- Order-preserving dictionary encoding
- Bit-level parallelization.
- Only require common instructions (no SIMD)

Implemented in Wisconsin's [QuickBit](#)

- Became an [Apache Incubator](#) project, but it then died in 2018.





# Bitweaving – Storage Layouts

---

## **Approach #1: Horizontal**

→ Row-oriented storage at the bit-level

## **Approach #2: Vertical**

→ Column-oriented storage at the bit-level

# Horizontal Storage

---

$t_0$ 

0	0	1
---	---	---

$t_1$ 

1	0	1
---	---	---

$t_2$ 

1	1	0
---	---	---

$t_3$ 

0	0	1
---	---	---

$t_4$ 

1	1	0
---	---	---

$t_5$ 

1	0	0
---	---	---

$t_6$ 

0	0	0
---	---	---

$t_7$ 

1	1	1
---	---	---

$t_8$ 

1	0	0
---	---	---

$t_9$ 

0	1	1
---	---	---

# Horizontal Storage

---

$$t_0 \begin{array}{|c|c|c|} \hline 0 & 0 & 1 \\ \hline \end{array} = 1$$

$$t_1 \begin{array}{|c|c|c|} \hline 1 & 0 & 1 \\ \hline \end{array} = 5$$

$$t_2 \begin{array}{|c|c|c|} \hline 1 & 1 & 0 \\ \hline \end{array} = 6$$

$$t_3 \begin{array}{|c|c|c|} \hline 0 & 0 & 1 \\ \hline \end{array} = 1$$

$$t_4 \begin{array}{|c|c|c|} \hline 1 & 1 & 0 \\ \hline \end{array} = 6$$

$$t_5 \begin{array}{|c|c|c|} \hline 1 & 0 & 0 \\ \hline \end{array} = 4$$

$$t_6 \begin{array}{|c|c|c|} \hline 0 & 0 & 0 \\ \hline \end{array} = 0$$

$$t_7 \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline \end{array} = 7$$

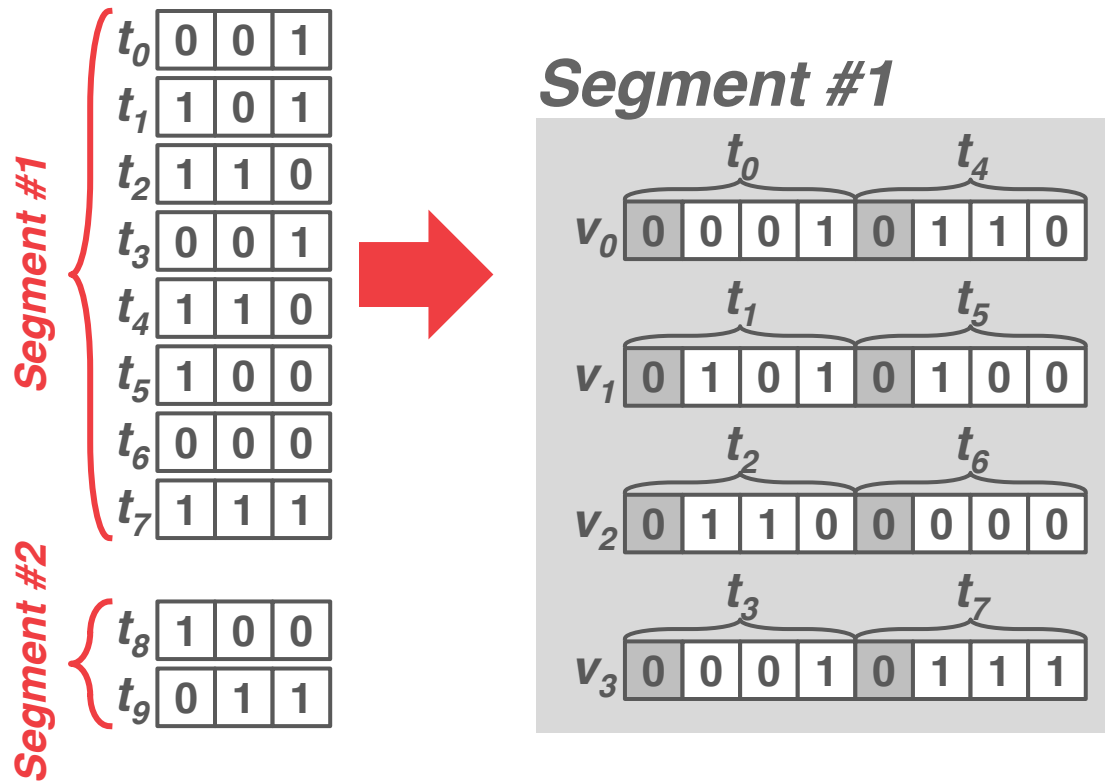
$$t_8 \begin{array}{|c|c|c|} \hline 1 & 0 & 0 \\ \hline \end{array} = 4$$

$$t_9 \begin{array}{|c|c|c|} \hline 0 & 1 & 1 \\ \hline \end{array} = 3$$


# Horizontal Storage

Segment #1	$t_0$	0	0	1	=1
	$t_1$	1	0	1	=5
	$t_2$	1	1	0	=6
	$t_3$	0	0	1	=1
	$t_4$	1	1	0	=6
	$t_5$	1	0	0	=4
	$t_6$	0	0	0	=0
	$t_7$	1	1	1	=7
Segment #2	$t_8$	1	0	0	=4
	$t_9$	0	1	1	=3


# Horizontal Storage




# Horizontal Storage




**Segment #1**




**Segment #2**



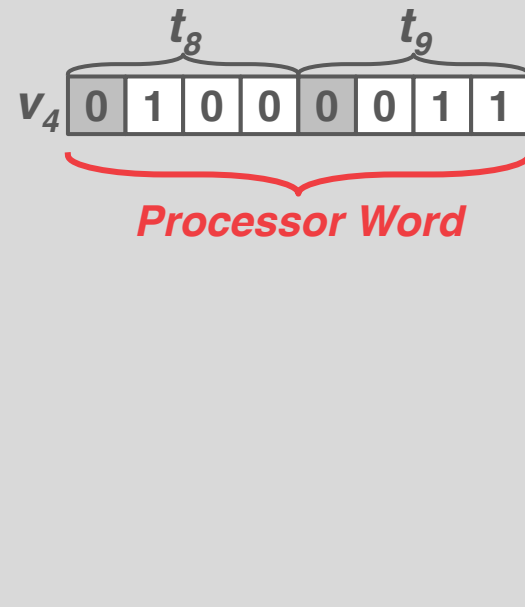
# Horizontal Storage



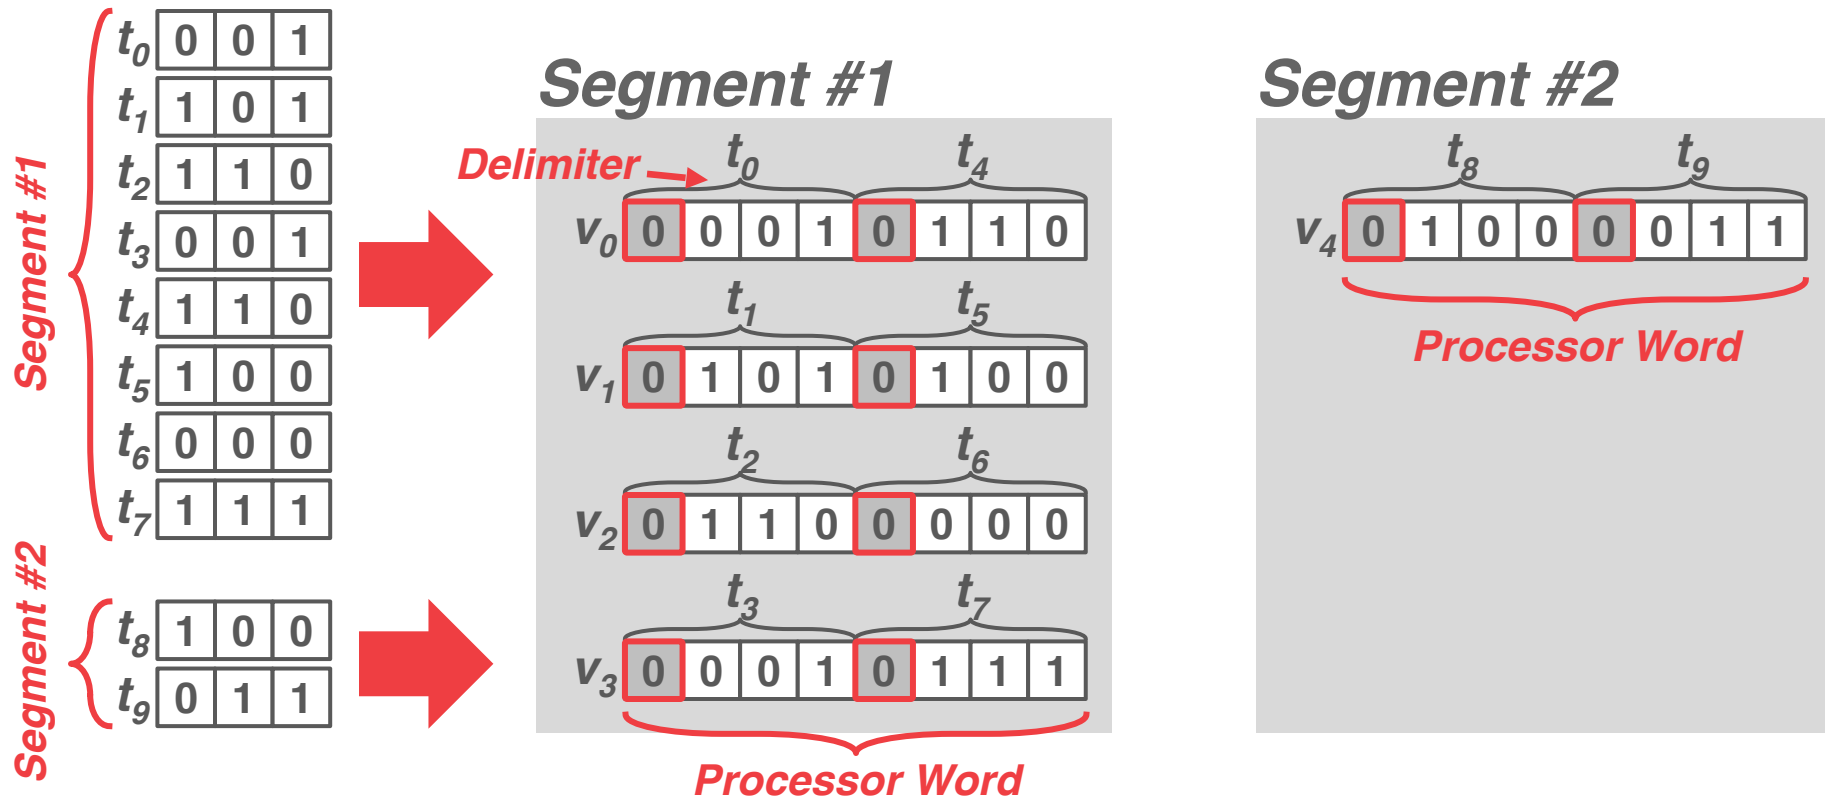
**Segment #1**



**Segment #2**



# Horizontal Storage





# Bitweaving/H – Example

---

```
SELECT * FROM table  
WHERE val < 5
```

$X =$ 


0	0	0	1	0	1	1	0
---	---	---	---	---	---	---	---

$Y =$ 

0	1	0	1	0	1	0	1
---	---	---	---	---	---	---	---


# Bitweaving/H – Example

```
SELECT * FROM table  
WHERE val < 5
```



# Bitweaving/H – Example

```
SELECT * FROM table  
WHERE val < 5
```



# Bitweaving/H – Example

SELECT \* FROM table  
WHERE val < 5

$$X = \begin{array}{|c|c|c|c|c|c|c|c|} \hline & \underbrace{\phantom{0001}}_{t_0} & & \underbrace{\phantom{0110}}_{t_4} & \\ \hline 0 & 0 & 0 & 1 & 0 & 1 & 1 & 0 \\ \hline \end{array}$$

$$Y = \begin{array}{|c|c|c|c|c|c|c|c|} \hline & \underbrace{\phantom{0101}}_5 & & \underbrace{\phantom{0101}}_5 & \\ \hline 0 & 1 & 0 & 1 & 0 & 1 & 0 & 1 \\ \hline \end{array}$$

$$mask = \begin{array}{|c|c|c|c|c|c|c|c|} \hline 0 & 1 & 1 & 1 & 0 & 1 & 1 & 1 \\ \hline \end{array}$$

$$(Y + (X \oplus mask)) \wedge \neg mask = \begin{array}{|c|c|c|c|c|c|c|c|} \hline 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ \hline \end{array}$$

# Bitweaving/H – Example

SELECT \* FROM table  
WHERE val < 5

$$X = \begin{array}{|c|c|c|c|c|c|c|c|} \hline & \underbrace{\phantom{0001}}_{t_0} & & \underbrace{\phantom{0110}}_{t_4} & \\ \hline 0 & 0 & 0 & 1 & 0 & 1 & 1 & 0 \\ \hline \end{array}$$

$$Y = \begin{array}{|c|c|c|c|c|c|c|c|} \hline & \underbrace{\phantom{0101}}_5 & & \underbrace{\phantom{0101}}_5 & \\ \hline 0 & 1 & 0 & 1 & 0 & 1 & 0 & 1 \\ \hline \end{array}$$

$$mask = \begin{array}{|c|c|c|c|c|c|c|c|} \hline 0 & 1 & 1 & 1 & 0 & 1 & 1 & 1 \\ \hline \end{array}$$

$$(Y + (X \oplus mask)) \wedge \neg mask = \begin{array}{|c|c|c|c|c|c|c|c|} \hline 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ \hline \end{array}$$

*Selection Vector*

# Bitweaving/H – Example

SELECT \* FROM table  
WHERE val < 5

$X =$ 

$t_0$				$t_4$			
0	0	0	1	0	1	1	0

$Y =$ 

5				5			
0	1	0	1	0	1	0	1

$mask =$ 

0	1	1	1	0	1	1	1
---	---	---	---	---	---	---	---

$(Y + (X \oplus mask)) \wedge \neg mask =$ 

1	0	0	0	0	0	0	0
1 < 5				5 < 6			

# Bitweaving/H – Example

SELECT \* FROM table  
WHERE val < 5

$X =$ 

$t_0$				$t_4$			
0	0	0	1	0	1	1	0

$Y =$ 

5				5			
0	1	0	1	0	1	0	1

$mask =$ 

0	1	1	1	0	1	1	1
---	---	---	---	---	---	---	---

$(Y + (X \oplus mask)) \wedge \neg mask =$ 

1	0	0	0	0	0	0	0
1 < 5				5 < 6			


Only requires four instructions to evaluate a single word.

Works on any word size and encoding length.

Paper contains algorithms for other operators.

# Bitweaving/H – Example


SELECT \* FROM table  
WHERE val < 5





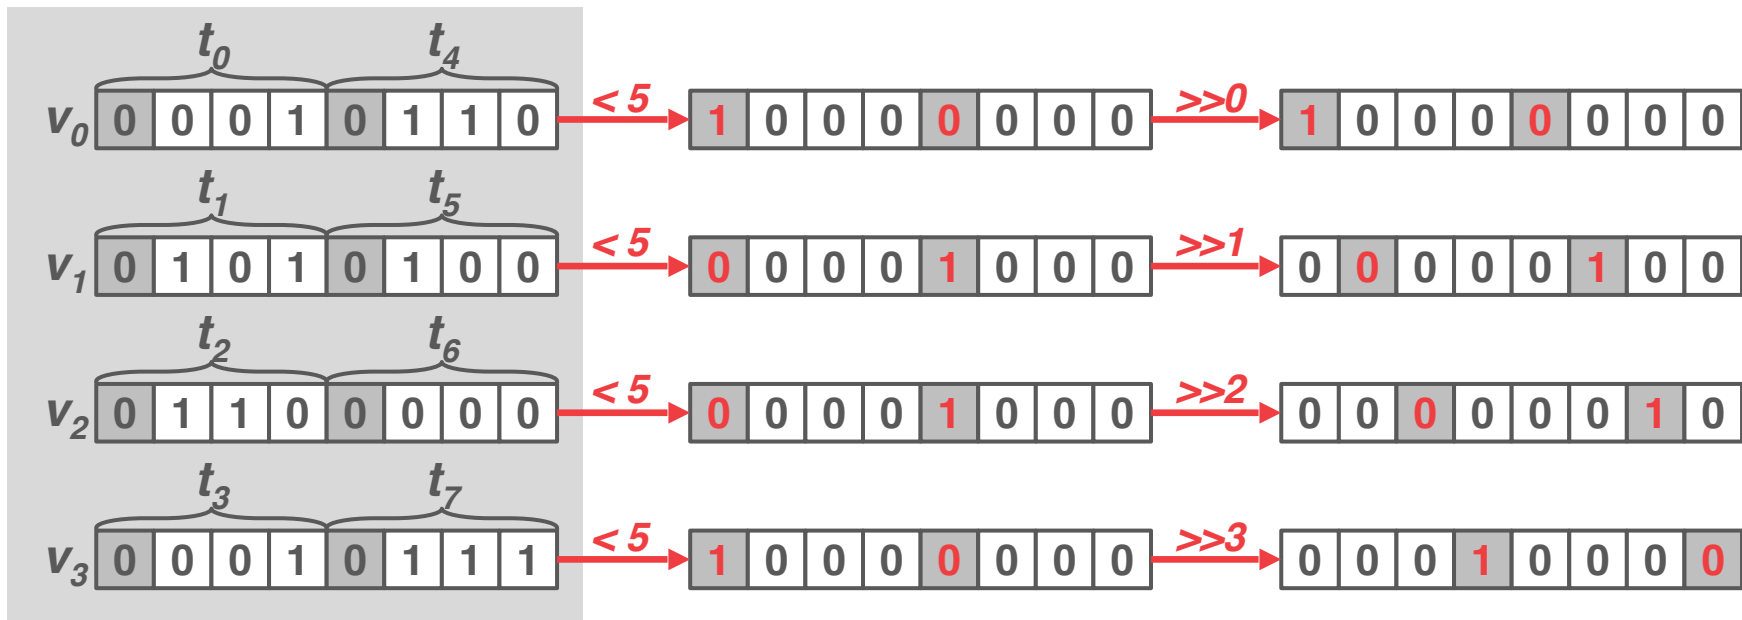
# Bitweaving/H – Example

SELECT \* FROM table  
WHERE val < 5



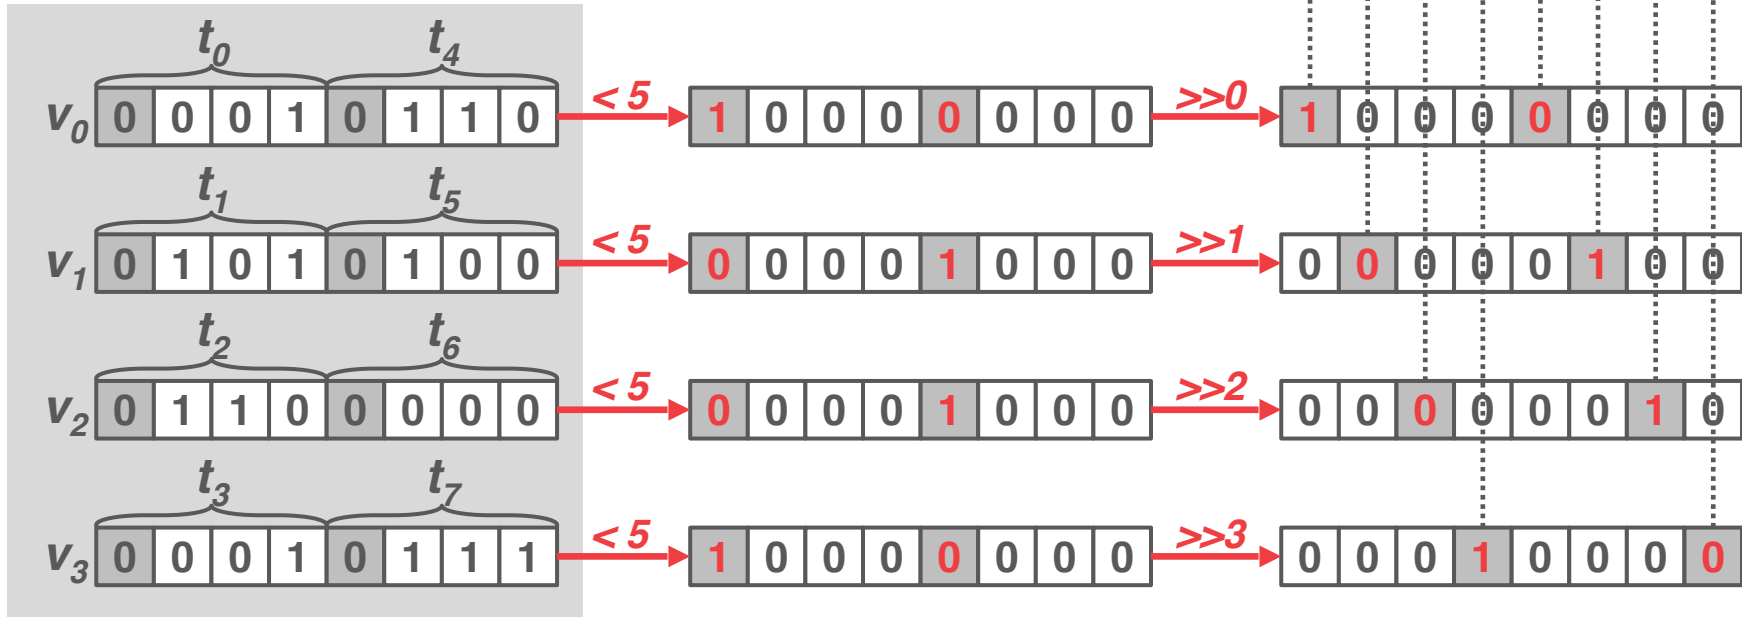
# Bitweaving/H – Example

SELECT \* FROM table  
WHERE val < 5



# Bitweaving/H – Example

SELECT \* FROM table  
WHERE val < 5



# Selection Vector

---

SIMD operators produce a bitmask specifying which tuples satisfy a predicate.

→ DBMS must convert it into column offsets.

**Approach #1: Iteration**

**Approach #2: Pre-computed Positions  
Table**

# Selection Vector

---

SIMD operators produce a bitmask specifying which tuples satisfy a predicate.

→ DBMS must convert it into column offsets.

**Approach #1: Iteration**

**Approach #2: Pre-computed Positions Table**

<i>Selection Vector</i>	$t_0$	$t_1$	$t_2$	$t_3$	$t_4$	$t_5$	$t_6$	$t_7$
	1	0	0	1	0	1	1	0

# Selection Vector

SIMD operators produce a bitmask specifying which tuples satisfy a predicate.

→ DBMS must convert it into column offsets.

**Approach #1: Iteration**

**Approach #2: Pre-computed Positions  
Table**

	$t_0$	$t_1$	$t_2$	$t_3$	$t_4$	$t_5$	$t_6$	$t_7$
<b>Selection Vector</b>	1	0	0	1	0	1	1	0

```
tuples = [ ]  
for (i=0; i<n; i++) {  
    if sv[i] == 1  
        tuples.add(i);  
}
```

# Selection Vector

---

SIMD operators produce a bitmask specifying which tuples satisfy a predicate.

→ DBMS must convert it into column offsets.

**Approach #1: Iteration**

**Approach #2: Pre-computed Positions Table**

<i>Selection Vector</i>	$t_0$	$t_1$	$t_2$	$t_3$	$t_4$	$t_5$	$t_6$	$t_7$
	1	0	0	1	0	1	1	0

→ DBMS must convert it into column offsets.

## Approach #1: Iteration

## Approach #2: Pre-computed Positions Table

**Selection Vector**

$t_0$	$t_1$	$t_2$	$t_3$	$t_4$	$t_5$	$t_6$	$t_7$
1	0	0	1	0	1	1	0

[illegible][illegible]




# Selection Vector

SIMD operators produce a bitmask specifying which tuples satisfy a predicate.

→ DBMS must convert it into column offsets.

## Approach #1: Iteration

## Approach #2: Pre-computed Positions Table



# Vertical Storage

---

Segment #1	$t_0$	0	0	1
	$t_1$	1	0	1
	$t_2$	1	1	0
	$t_3$	0	0	1
	$t_4$	1	1	0
	$t_5$	1	0	0
	$t_6$	0	0	0
	$t_7$	1	1	1
Segment #2	$t_8$	1	0	0
	$t_9$	0	1	1

# Vertical Storage

---

Segment #1	$t_0$	0	0	1
	$t_1$	1	0	1
	$t_2$	1	1	0
	$t_3$	0	0	1
	$t_4$	1	1	0
	$t_5$	1	0	0
	$t_6$	0	0	0
	$t_7$	1	1	1
Segment #2	$t_8$	1	0	0
	$t_9$	0	1	1

# Vertical Storage

**Segment #1**

$t_0$	0	0	1
$t_1$	1	0	1
$t_2$	1	1	0
$t_3$	0	0	1
$t_4$	1	1	0
$t_5$	1	0	0
$t_6$	0	0	0
$t_7$	1	1	1



**Segment #1**

	$t_0$	$t_1$	$t_2$	$t_3$	$t_4$	$t_5$	$t_6$	$t_7$
$v_0$	0	1	1	0	1	1	0	1
$v_1$	0	0	1	0	1	0	0	1
$v_2$	1	1	0	1	0	0	0	1

**Segment #2**

$t_8$	1	0	0
$t_9$	0	1	1




**Segment #2**

	$t_8$	$t_9$	-	-	-	-	-
$v_3$	1	0	0	0	0	0	0
$v_4$	0	1	0	0	0	0	0
$v_5$	0	1	0	0	0	0	0

# Vertical Storage

*Segment #1*

$t_0$	0	0	1
$t_1$	1	0	1
$t_2$	1	1	0
$t_3$	0	0	1
$t_4$	1	1	0
$t_5$	1	0	0
$t_6$	0	0	0
$t_7$	1	1	1



*Segment #1*


	$t_0$	$t_1$	$t_2$	$t_3$	$t_4$	$t_5$	$t_6$	$t_7$
$v_0$	0	1	1	0	1	1	0	1
$v_1$	0	0	1	0	1	0	0	1
$v_2$	1	1	0	1	0	0	0	1

*Segment #2*


	$t_8$	$t_9$	-	-	-	-	-	-
$v_3$	1	0	0	0	0	0	0	0
$v_4$	0	1	0	0	0	0	0	0
$v_5$	0	1	0	0	0	0	0	0

*Segment #2*

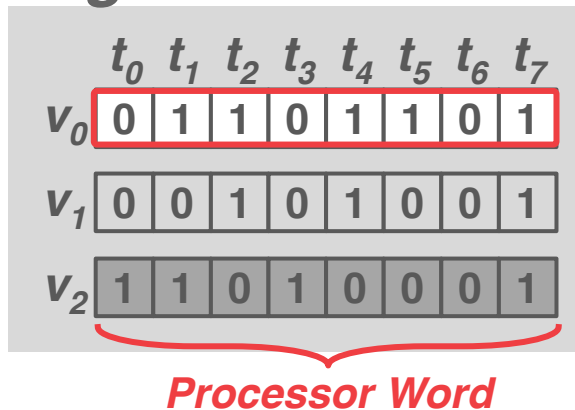
$t_8$	1	0	0
$t_9$	0	1	1




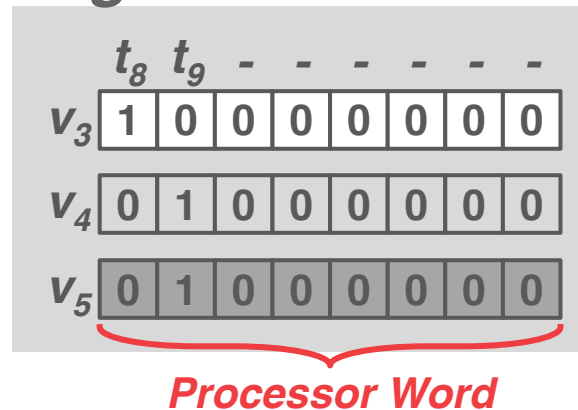
# Vertical Storage



**Segment #1**



**Segment #2**



# Bitweaving/V – Example


```
SELECT * FROM table  
WHERE val = 2
```

## Segment #1

	$t_0$	$t_1$	$t_2$	$t_3$	$t_4$	$t_5$	$t_6$	$t_7$
$v_0$	0	1	1	0	1	1	0	1
$v_1$	0	0	1	0	1	0	0	1
$v_2$	1	1	0	1	0	0	0	1

# Bitweaving/V – Example

SELECT \* FROM table  
WHERE val = 2




## Segment #1

	$t_0$	$t_1$	$t_2$	$t_3$	$t_4$	$t_5$	$t_6$	$t_7$
$v_0$	0	1	1	0	1	1	0	1
$v_1$	0	0	1	0	1	0	0	1
$v_2$	1	1	0	1	0	0	0	1



# Bitweaving/V – Example

SELECT \* FROM table  
WHERE val = 2



## Segment #1

	$t_0$	$t_1$	$t_2$	$t_3$	$t_4$	$t_5$	$t_6$	$t_7$
$v_0$	0	1	1	0	1	1	0	1
$v_1$	0	0	1	0	1	0	0	1
$v_2$	1	1	0	1	0	0	0	1

# Bitweaving/V – Example

SELECT \* FROM table  
WHERE val = 2

0 1 0

*Segment #1*

	$t_0$	$t_1$	$t_2$	$t_3$	$t_4$	$t_5$	$t_6$	$t_7$
$v_0$	0	1	1	0	1	1	0	1
$v_1$	0	0	1	0	1	0	0	1
$v_2$	1	1	0	1	0	0	0	1


0 0 0 0 0 0 0 0

SIMDCompare

1 0 0 1 0 0 1 0

# Bitweaving/V – Example

SELECT \* FROM table  
WHERE val = 2



*Segment #1*

	$t_0$	$t_1$	$t_2$	$t_3$	$t_4$	$t_5$	$t_6$	$t_7$
$v_0$	0	1	1	0	1	1	0	1
$v_1$	0	0	1	0	1	0	0	1
$v_2$	1	1	0	1	0	0	0	1

SIMDCompare



# Bitweaving/V – Example

SELECT \* FROM table  
WHERE val = 2

0 1 0

*Segment #1*

	$t_0$	$t_1$	$t_2$	$t_3$	$t_4$	$t_5$	$t_6$	$t_7$
$v_0$	0	1	1	0	1	1	0	1
$v_1$	0	0	1	0	1	0	0	1
$v_2$	1	1	0	1	0	0	0	1

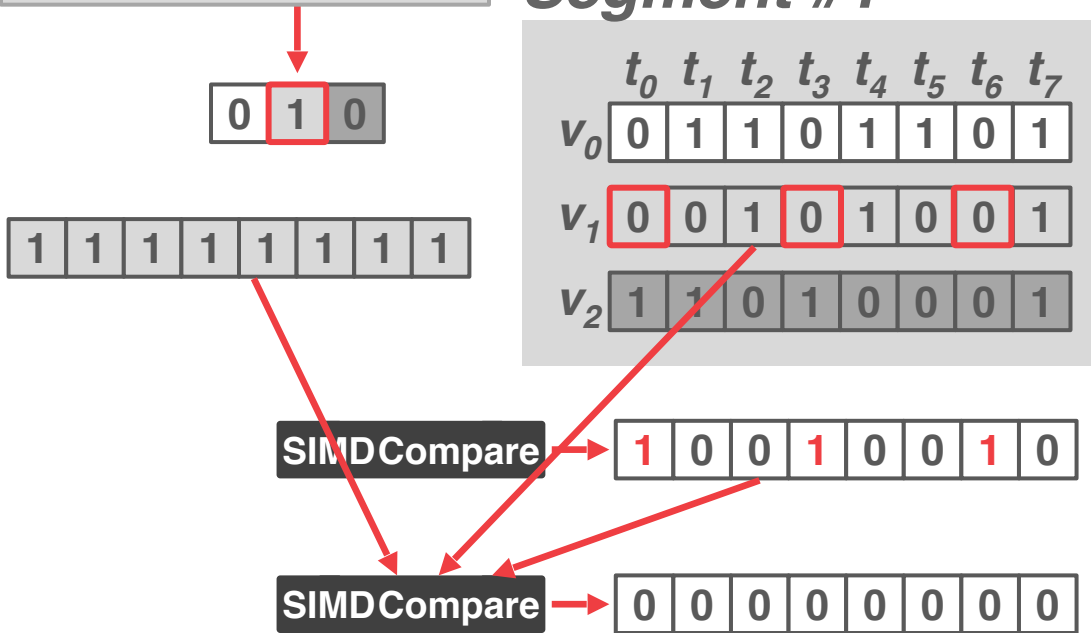
1 1 1 1 1 1 1 1

SIMDCompare

1 0 0 1 0 0 1 0


SIMDCompare

0 0 0 0 0 0 0 0



# Bitweaving/V – Example

SELECT \* FROM table  
WHERE val = 2



## Segment #1

	$t_0$	$t_1$	$t_2$	$t_3$	$t_4$	$t_5$	$t_6$	$t_7$
$v_0$	0	1	1	0	1	1	0	1
$v_1$	0	0	1	0	1	0	0	1
$v_2$	1	1	0	1	0	0	0	1



SIMDCompare



SIMDCompare



Can perform early pruning just like in BitMap indexes.

The last vector is skipped because all bits in previous comparison are zero.

# Today's Agenda

---

~~Zone Maps~~

~~Bitmap Indexes~~

~~Bit-Slicing~~

~~Bit-Weaving~~

Column Imprints

Column Sketches

# Observation

---

All the previous Bitmap schemes were about storing exact/lossless representations of columnar data.

The DBMS could trade off accuracy for faster evaluation in the common case.

→ It still must always check the original data to avoid false positives.

# Column Imprints

---

Store a bitmap that indicates whether there is a bit set at a bit-slice of cache-line values.





# Column Imprints

---

Store a bitmap that indicates whether there is a bit set at a bit-slice of cache-line values.

## *Original Data*

value
1
8
4



# Column Imprints

Store a bitmap that indicates whether there is a bit set at a bit-slice of cache-line values.

*Original Data*

value
1
8
4



*Bitmap Indexes*

1	2	3	4	5	6	7	8
1	0	0	0	0	0	0	0
0	0	0	0	0	0	0	1
0	0	0	1	0	0	0	0



# Column Imprints

Store a bitmap that indicates whether there is a bit set at a bit-slice of cache-line values.

*Original Data*

value
1
8
4



*Bitmap Indexes*

1	2	3	4	5	6	7	8
1	0	0	0	0	0	0	0
0	0	0	0	0	0	0	1
0	0	0	1	0	0	0	0



*Column Imprint*

1	0	0	1	0	0	0	1
---	---	---	---	---	---	---	---



# Column Sketches

---

A variation of range-encoded Bitmaps that uses smaller “sketch” codes to indicate that a tuple's value exists in a range.

DBMS must automatically figure out the best mapping of codes.

- Trade-off between distribution of values and compactness.
- Assign unique codes to frequent values to avoid false positives.

# Column Sketches

---

## *Original Data*

val
13
191
56
92
81
120
231
172

*8-bits*

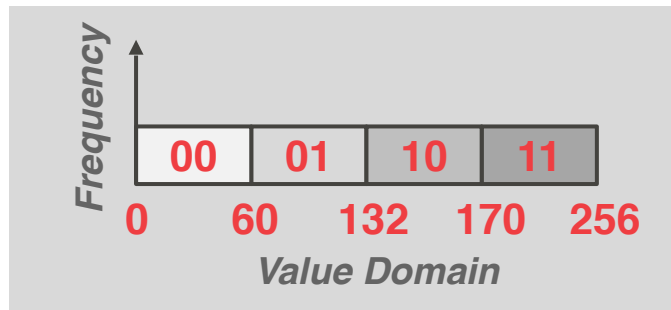
# Column Sketches

*Original Data*

val
13
191
56
92
81
120
231
172

*8-bits*

*Histogram*




# Column Sketches

*Original Data*

val
13
191
56
92
81
120
231
172

8-bits

*Histogram*



*Compression Map*

60	→	00
132	→	01
170	→	10
256	→	11


# Column Sketches

*Original Data*

val
13
191
56
92
81
120
231
172

*8-bits*

*Histogram*



*Compression Map*

60	→	00
132	→	01
170	→	10
256	→	11

*Sketched Column*


code
00
10
11
00
01
01
10
11

*2-bits*




# Column Sketches


## Histogram



## Sketched Column




## Compression Map




SELECT \* FROM table  
WHERE val < 90

# Column Sketches

## Histogram



## Sketched Column




## Compression Map

60	→	00
132	→	01
170	→	10
256	→	11


SELECT \* FROM table  
WHERE val < 90

# Column Sketches

## Histogram



## Sketched Column



## Compression Map

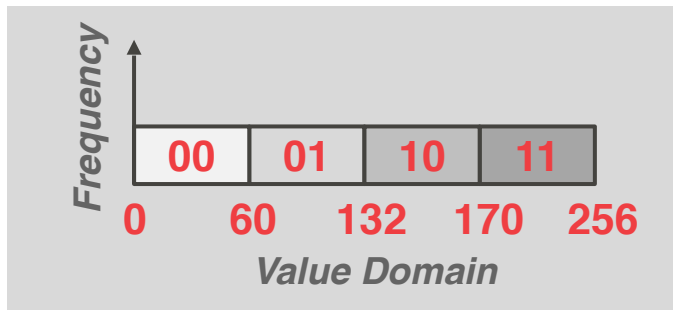
60	→	00
132	→	01
170	→	10
256	→	11

map(90) = 01

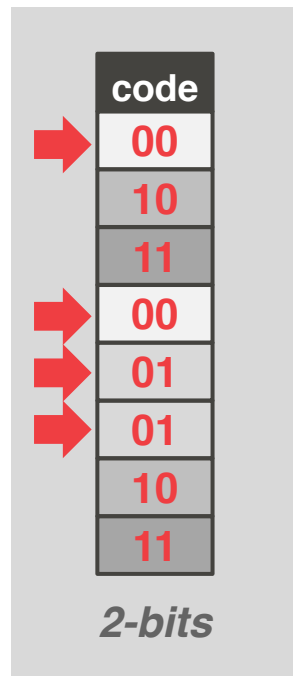
SELECT \* FROM table  
WHERE val < 90

# Column Sketches

## Histogram



## Sketched Column



## Compression Map

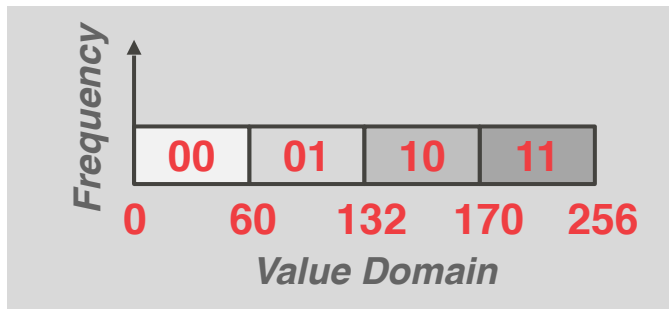
60	→	00
132	→	01
170	→	10
256	→	11

map(90) = 01

SELECT \* FROM table  
WHERE val < 90

# Column Sketches

## Histogram




## Compression Map

60	→	00
132	→	01
170	→	10
256	→	11

map(90) = 01

SELECT \* FROM table  
WHERE val < 90

## Sketched Column



# Parting Thoughts

---

Zone Maps are the most widely used method to accelerate sequential scans.

Bitmap indexes are more common in NSM DBMSs than columnar OLAP systems.

We're ignoring multi-dimensional and inverted indexes...

# Next Class

---

Data Compression (Tuples, Indexes)